

Homework Assignment 3

1. We perform best subset, forward stepwise, and backward stepwise selection on a single data set. For each approach, we obtain $p + 1$ models, containing $0, 1, 2, \dots, p$ predictors. Explain your answers:

(a) Which of the three models with k predictors has the smallest training RSS?

Given the condition that a single data set is used to perform the best subset, forward stepwise, and backward stepwise selection and for each approach, we obtain $p + 1$ models, containing $0, 1, 2, \dots, p$ predictors then the out of the three models with k predictors best subset selection method will have the smallest training RSS. Though forward stepwise tends to do well in practice, it is not guaranteed to find the best possible model out of all the models containing subsets of the p predictors. Also like forward stepwise selection, backward stepwise selection is not guaranteed to yield the best model containing a subset of the p predictors. Furthermore, Backward selection requires that the number of samples n is larger than the number of variables p (so that the full model can be fit). In contrast, forward stepwise can be used even when $n < p$, and so is the only viable subset method when p is very large.

(b) Which of the three models with k predictors has the smallest test RSS?

It might be possible that the best subset selection method and backward stepwise selection might arrive at the same model as backward stepwise selection iteratively removes the least useful predictor, one at a time and best subset selection tries different combinations of all the predictors while forward selection retains the old predictors when it iteratively add new predictors. So there is no guarantee that which model will come up with the smallest test RSS as either of these methods can pick up a model that best fits the test data and provide the smallest test RSS.

(c) True or False:

- i. The predictors in the k -variable model identified by forward stepwise are a subset of the predictors in the $(k+1)$ -variable model identified by forward stepwise selection. - **True**
- ii. The predictors in the k -variable model identified by backward stepwise are a subset of the predictors in the $(k + 1)$ -variable model identified by backward stepwise selection. - **True**
- iii. The predictors in the k -variable model identified by backward stepwise are a subset of the predictors in the $(k + 1)$ -variable model identified by forward stepwise selection. - **False**
- iv. The predictors in the k -variable model identified by forward stepwise are a subset of the predictors in the $(k+1)$ -variable model identified by backward stepwise selection. - **False**
- v. The predictors in the k -variable model identified by best subset are a subset of the predictors in the $(k + 1)$ -variable model identified by best subset selection. - **False**

2. We will now explore (6.12) and (6.13) further.

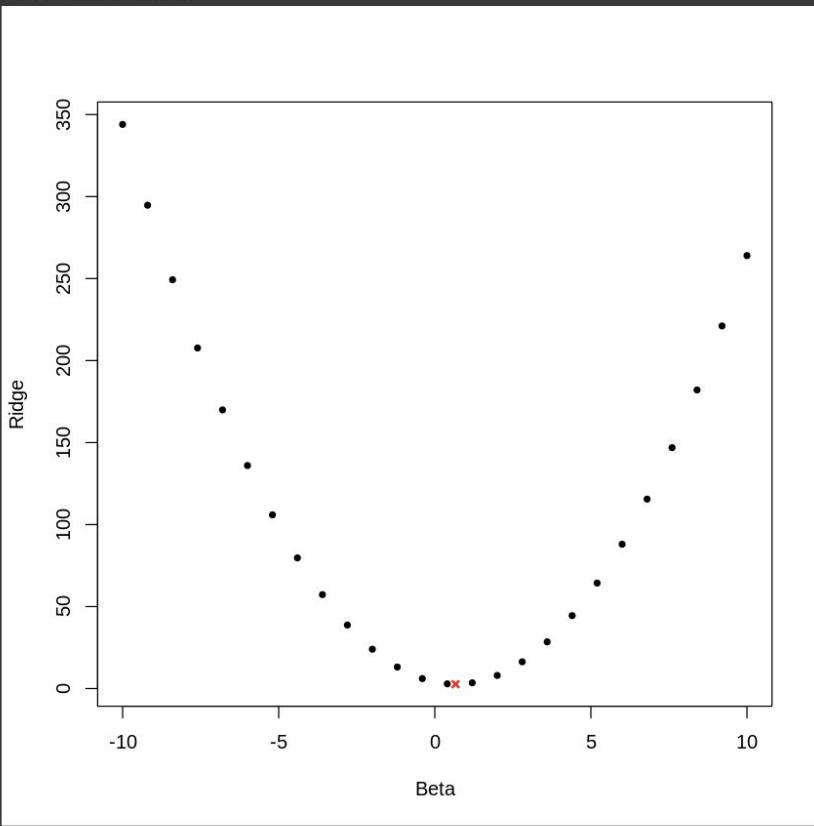
(a) Consider (6.12) with $p = 1$. For some choice of y_1 and $\lambda > 0$, plot (6.12) as a function of β_1 . Your plot should confirm that (6.12) is solved by (6.14).

$$\sum_{j=1}^p (y_j - \beta_j)^2 + \lambda \sum_{j=1}^p \beta_j^2 \quad (6.12)$$

$$\hat{\beta}_j^R = y_j / (1 + \lambda), \quad (6.14)$$

```
1 betas=seq(-10,10,0.8)
2
3 # equation 6.12
4 func1 = ((y - betas)^2) + (lambda * (betas^2))
5
6 plot(betas,func1,pch=20,xlab='Beta',ylab='Ridge')
7
8 # equation 6.14
9 estimation_beta=y/(1+lambda)
10 print(estimation_beta)
11
12 estimation_function=((y - estimation_beta)^2) + (lambda * (estimation_beta^2))
13 print(estimation_function)
14
15 points(estimation_beta,estimation_function,col='red',pch=4,lwd=2,cex=estimation_beta)

[1] 0.6666667
[1] 2.666667
```



The red cross in the figure shows that the function (6.12) is minimized by beta (6.14)

(b) Consider (6.13) with $p = 1$. For some choice of y_1 and $\lambda > 0$, plot (6.13) as a function of β_1 . Your plot should confirm that (6.13) is solved by (6.15).

$$\sum_{j=1}^p (y_j - \beta_j)^2 + \lambda \sum_{j=1}^p |\beta_j| \quad (6.13)$$

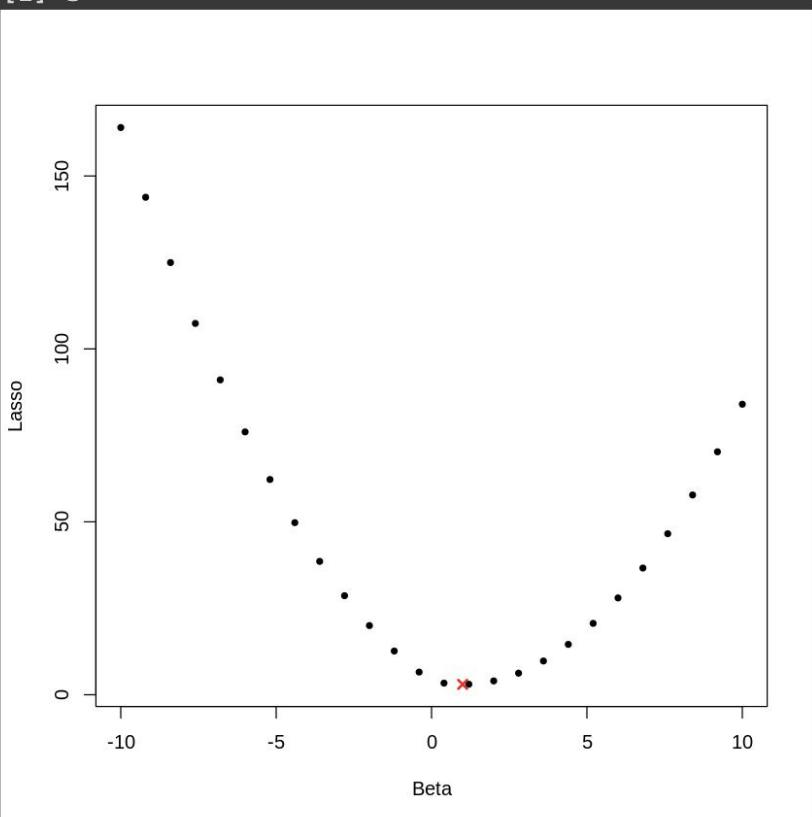
$$\hat{\beta}_j^L = \begin{cases} y_j - \lambda/2 & \text{if } y_j > \lambda/2; \\ y_j + \lambda/2 & \text{if } y_j < -\lambda/2; \\ 0 & \text{if } |y_j| \leq \lambda/2. \end{cases} \quad (6.15)$$

```

1 betas=seq(-10,10,0.8)
2
3 # equation 6.13
4 func1=(y - betas)^2 + (lambda * abs(betas))
5
6 plot(betas,func1,pch=20,xlab='Beta',ylab='Lasso')
7
8 # equation 6.15
9 # since y > (lambda / 2) i.e 2 > (2/2) so
10 estimation_beta=y-(lambda/2)
11 print(estimation_beta)
12
13 estimation_function=(y - estimation_beta)^2 + (lambda * abs(estimation_beta))
14 print(estimation_function)
15
16 points(estimation_beta,estimation_function,col='red',pch=4,lwd=2,cex=estimation_beta)

[1] 1
[1] 3

```



The red cross in the figure shows that the function (6.13) is minimized by beta (6.15)

3. In this exercise, we will generate simulated data, and will then use this data to perform best subset selection.

(a) Use the rnorm() function to generate a predictor X of length n = 100, as well as a noise vector ϵ of length n = 100.

```
1 set.seed(123)
2
3 x_n=100
4 e_n=100
5 x=rnorm(x_n)
6 print(x)
7 e_noise=rnorm(e_n)
8 print(e_noise)

[1] -0.560475647 -0.230177489  1.558708314  0.070508391  0.129287735
[6]  1.715064987  0.460916206 -1.265061235 -0.686852852 -0.445661970
[11] 1.224081797  0.359813827  0.400771451  0.110682716 -0.555841135
[16] 1.786913137  0.497850478 -1.966617157  0.701355902 -0.472791408
[21] -1.067823706 -0.217974915 -1.026004448 -0.728891229 -0.625039268
[26] -1.686693311  0.837787044  0.153373118 -1.138136937  1.253814921
[31] 0.426464221 -0.295071483  0.895125661  0.878133488  0.821581082
[36] 0.688640254  0.553917654 -0.061911711 -0.305962664 -0.380471001
[41] -0.694706979 -0.207917278 -1.265396352  2.168955965  1.207961998
[46] -1.123108583 -0.402884835 -0.466655354  0.779965118 -0.083369066
[51] 0.253318514 -0.028546755 -0.042870457  1.368602284 -0.225770986
[56] 1.516470604 -1.548752804  0.584613750  0.123854244  0.215941569
[61] 0.379639483 -0.502323453 -0.333207384 -1.018575383 -1.071791226
[66] 0.303528641  0.448209779  0.053004227  0.922267468  2.050084686
[71] -0.491031166 -2.309168876  1.005738524 -0.709200763 -0.688008616
[76] 1.025571370 -0.284773007 -1.220717712  0.181303480 -0.138891362
[81] 0.005764186  0.385280401 -0.370660032  0.644376549 -0.220486562
[86] 0.331781964  1.096839013  0.435181491 -0.325931586  1.148807618
[91] 0.993503856  0.548396960  0.238731735 -0.627906076  1.360652449
[96] -0.600259587  2.187332993  1.532610626 -0.235700359 -1.026420900
[1] -0.71040656  0.25688371 -0.24669188 -0.34754260 -0.95161857 -0.04502772
[7] -0.78490447 -1.66794194 -0.38022652  0.91899661 -0.57534696  0.60796432
[13] -1.61788271 -0.05556197  0.51940720  0.30115336  0.10567619 -0.64070601
[19] -0.84970435 -1.02412879  0.11764660 -0.94747461 -0.49055744 -0.25609219
[25] 1.84386201 -0.65194990  0.23538657  0.07796085 -0.96185663 -0.07130809
[31] 1.44455086  0.45150405  0.04123292 -0.42249683 -2.05324722  1.13133721
[37] -1.46064007  0.73994751  1.90910357 -1.44389316  0.70178434 -0.26219749
[43] -1.57214416 -1.51466765 -1.60153617 -0.53090652 -1.46175558  0.68791677
[49] 2.10010894 -1.28703048  0.78773885  0.76904224  0.33220258 -1.00837661
[55] -0.11945261 -0.28039534  0.56298953 -0.37243876  0.97697339 -0.37458086
[61] 1.05271147 -1.04917701 -1.26015524  3.24103993 -0.41685759  0.29822759
[67] 0.63656967 -0.48378063  0.51686204  0.36896453 -0.21538051  0.06529303
[73] -0.03406725  2.12845190 -0.74133610 -1.09599627  0.03778840  0.31048075
[79] 0.43652348 -0.45836533 -1.06332613  1.26318518 -0.34965039 -0.86551286
[85] -0.23627957 -0.19717589  1.10992029  0.08473729  0.75405379 -0.49929202
[91] 0.21444531 -0.32468591  0.09458353 -0.89536336 -1.31080153  1.99721338
[97] 0.60070882 -1.25127136 -0.61116592 -1.18548008
```

(b) Generate a response vector Y of length n = 100 according to the model $Y = \beta_0 + \beta_1 X + \beta_2 X^2 + \beta_3 X^3 + \epsilon$, where $\beta_0, \beta_1, \beta_2$, and β_3 are constants of your choice.

```
1 beta_0=1
2 beta_1=2
3 beta_2=3
4 beta_3=4

1 y = beta_0 + (beta_1 * x) + (beta_2 * (x^2)) + (beta_3 * (x^3)) + e_noise

1 y

-0.593214479535489 · 0.906693002891148 · 26.3074134401957 · 0.809790594284573 · 0.365747201112343 · 33.3885432561223 · 2.16593425634838 · -6.49523905059849 · -0.634769301967035 · 1.26945656292057 · 14.7044857305433 ·
2.90232456031134 · 0.922997503362667 · 1.20797920407416 · 0.647673734916369 · 37.2770284501246 · 3.33852156425047 · -22.39541184983 · 4.40869591337405 · -0.721851957642085 · -2.46759151777556 · -0.282311877017288 ·
-2.7047494522353 · -0.669015775289461 · 1.78905914898594 · -13.6846603282348 · 7.36674984033212 · 1.46970840065668 · -4.24923206353372 · 16.0367256213954 · 4.1533416447014 · 1.0197984595414 · 8.10411165252995 ·
7.35570504509802 · 4.83315540839494 · 6.23757671894972 · -2.24749215656729 · 1.62667402426127 · -2.4634491796417 · -0.99086577733988 · 0.419111959316267 · 0.415704019583014 · -6.40400498779967 · 58.7506395370841 ·
13.2424026059506 · -3.65963989559314 · -1.04215561077853 · 1.00141876844334 · 8.38302928200068 · -0.4352351990908 · 2.55190875334469 · 1.71430042905389 · 1.25166013033574 · 18.6020082240071 · 0.535890555395325 ·
24.6012024028037 · -9.19818259985384 · 3.62132978466808 · 2.27830112937003 · 1.23747264188042 · 3.46323272992399 · -0.80384015835976 · -0.741468809209702 · 1.08930582900584 · -3.03905326399369 · 2.29352971598538 ·
3.49583229380926 · 0.631251822392155 · 9.05096785256671 · 52.5424464259703 · 0.0523187261096945 · -36.808625622539 · 10.0811979513454 · 1.79213284541305 · -0.99977391458018 · 9.42530643503964 · 0.619153955559616 ·
-3.936717933768 · 1.92158176546959 · 0.311007066246303 · -0.0516973185744989 · 3.70783455363648 · 0.117497890337623 · 3.73913869390281 · 0.425715058684329 · 1.94271511586668 · 13.1910000151638 · 2.8529128469592 ·
1.28238813694678 · 12.8221965708793 · 10.0851543305882 · 3.33402358993131 · 1.79744952206143 · -0.958625558264435 · 18.0409407947742 · 2.01250702461877 · 62.1891797422247 · 24.2604036646214 · 0.031720303908158 ·
-3.4032036367233
```

(c) Use the regsubsets() function to perform best subset selection in order to choose the best model containing the predictors X, X^2, \dots, X^{10} . What is the best model obtained according to Cp, BIC, and adjusted R²? Show some plots to provide evidence for your answer, and report the coefficients of the best model obtained. Note you will need to use the data.frame() function to create a single data set containing both X and Y.

```
1 library(leaps)

1 question_3_data=data.frame(x,x^2,x^3,x^4,x^5,x^6,x^7,x^8,x^9,x^10,y)

1 model_bss=regsubsets(y~.,data=question_3_data,nvmax=10)

1 model_bss_summary=summary(model_bss)

1 model_bss_summary

Subset selection object
Call: regsubsets.formula(y ~ ., data = question_3_data, nvmax = 10)
10 Variables (and intercept)
      Forced in      Forced out
x        FALSE        FALSE
x.2      FALSE        FALSE
x.3      FALSE        FALSE
x.4      FALSE        FALSE
x.5      FALSE        FALSE
x.6      FALSE        FALSE
x.7      FALSE        FALSE
x.8      FALSE        FALSE
x.9      FALSE        FALSE
x.10     FALSE        FALSE
1 subsets of each size up to 10
Selection Algorithm: exhaustive
              x   x.2 x.3 x.4 x.5 x.6 x.7 x.8 x.9 x.10
1 ( 1 )   " " " " "*" " " " " " " " "
2 ( 1 )   " " "*" "*" " " " " " " " " "
3 ( 1 )   "*" "*" "*" " " " " " " " " "
4 ( 1 )   "*" "*" "*" " " " " " "*" " " " "
5 ( 1 )   "*" "*" "*" " " " " " "*" " " "*" " "
6 ( 1 )   "*" "*" "*" "*" " " "*" " " "*" " "
7 ( 1 )   "*" "*" "*" "*" " " "*" " " "*" " "
8 ( 1 )   "*" "*" "*" "*" "*" " " "*" " " "*" " "
9 ( 1 )   "*" "*" "*" "*" "*" "*" " " "*" " "
10 ( 1 )  "*" "*" "*" "*" "*" "*" "*" " " "*" " "

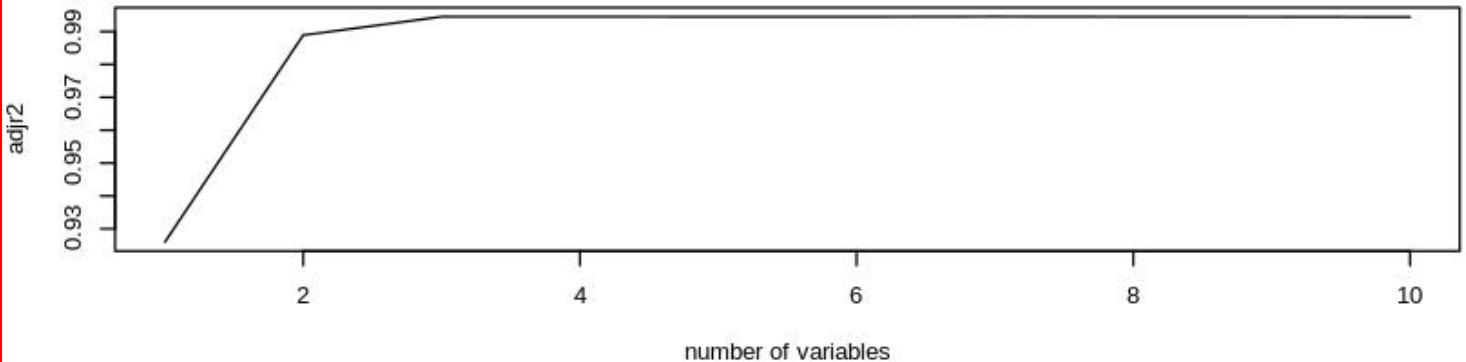
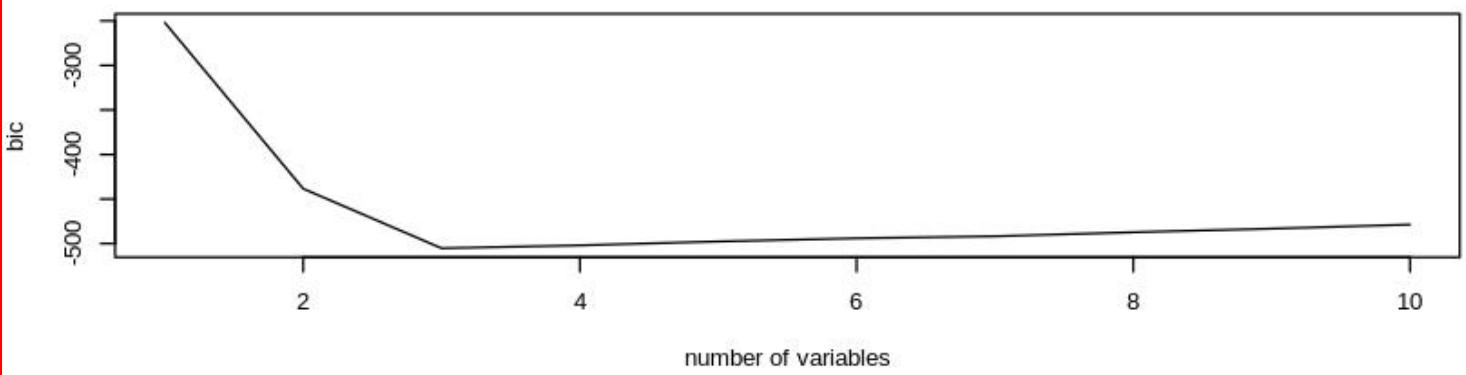
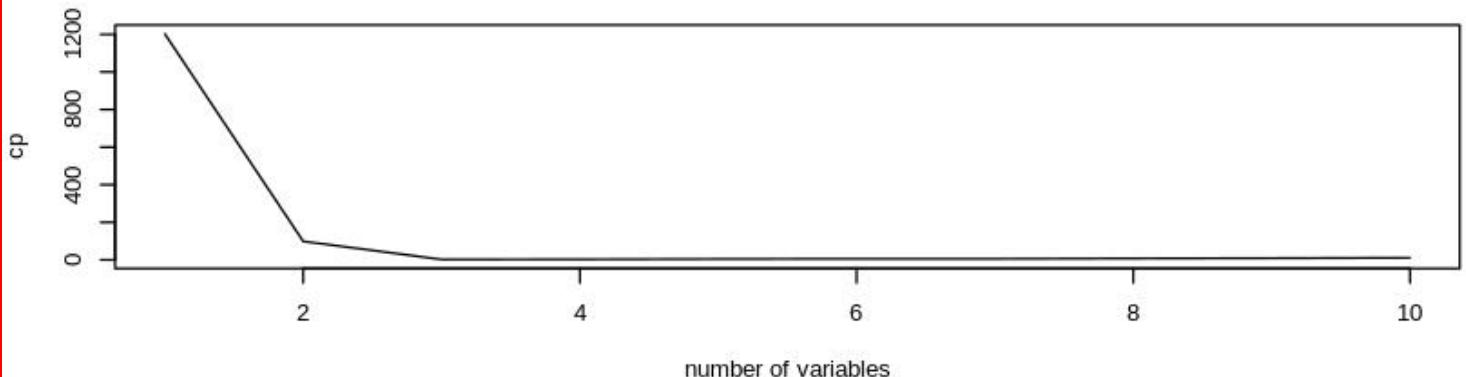
1 names(model_bss)

'np' · 'nrbar' · 'd' · 'rbar' · 'thetab' · 'first' · 'last' · 'vorder' · 'tol' · 'rss' · 'bound' · 'nvmax' · 'ress' · 'ir' · 'nbest'
```

```

par(mfrow=c(3,1))
plot(model_bss_summary$cp,xlab='number of variables',ylab='cp',type='l')
plot(model_bss_summary$bic,xlab='number of variables',ylab='bic',type='l')
plot(model_bss_summary$adjr2,xlab='number of variables',ylab='adjr2',type='l')

```



```
1 c(which.min(model_bss_summary$cp),which.min(model_bss_summary$bic),which.max(model_bss_summary$adjr2))
```

3 3 7

```
1 print(coef(model_bss,3))
```

	x	x.2	x.3
(Intercept)	0.9703939	1.9204462	2.9084569
			4.0204363

```
1 print(coef(model_bss,7))
```

	x	x.2	x.3	x.4	x.6
(Intercept)	0.7931750	1.8907484	5.2758848	4.0530514	-4.0243410
x.8		x.10			2.2138377

(d) Repeat (c), using forward stepwise selection and also using back-wards stepwise selection. How does your answer compare to the results in (c)?

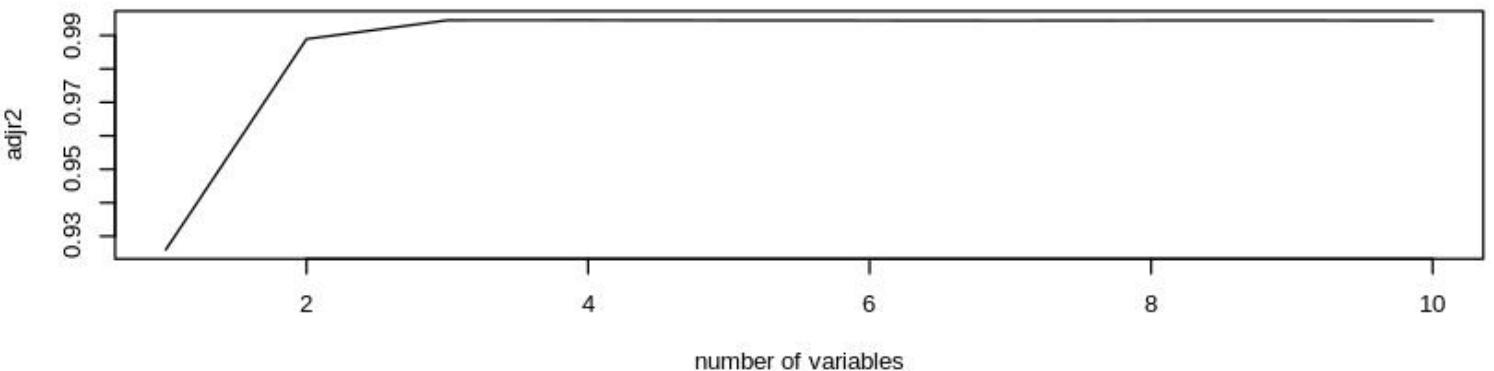
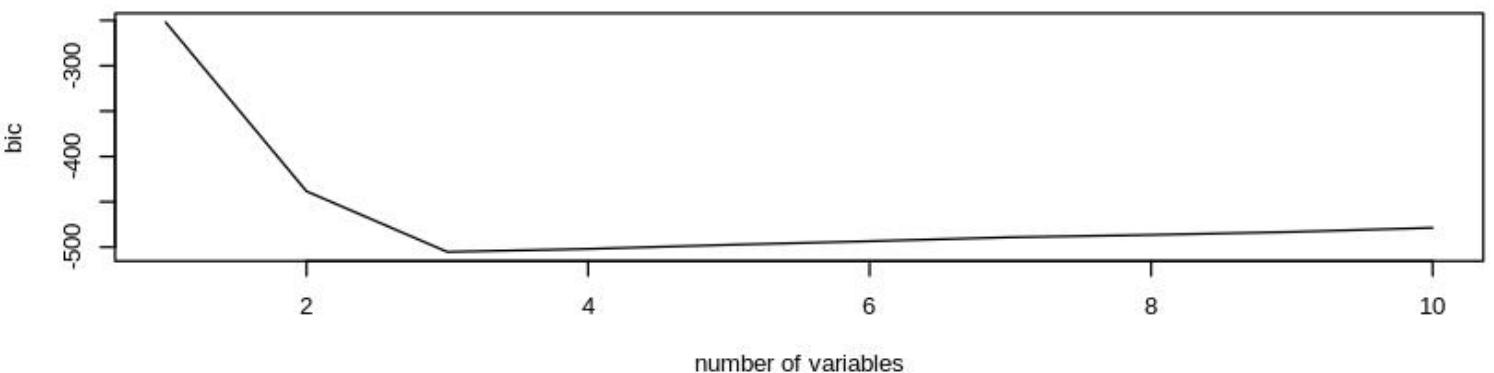
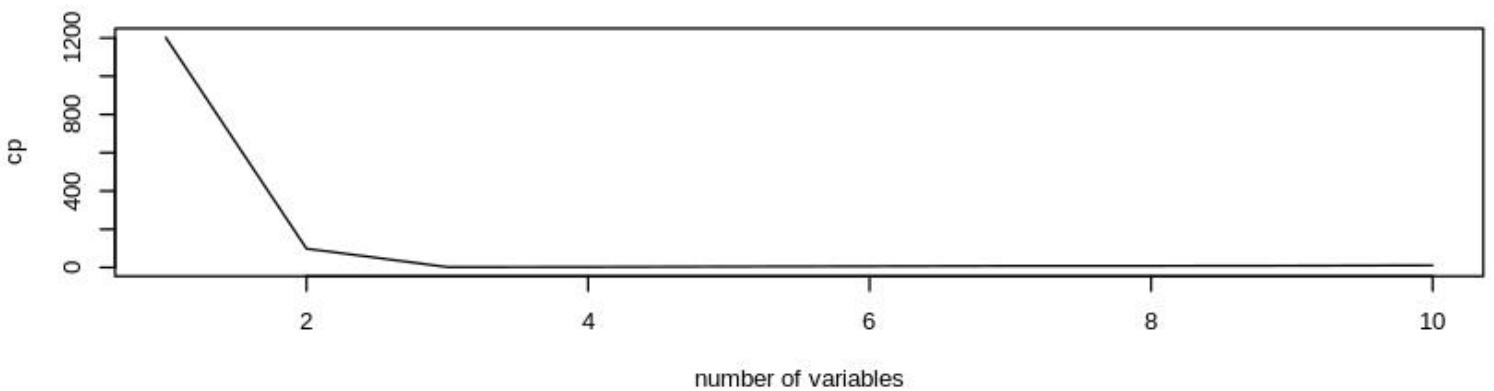
Forward Stepwise Selection

```
1 # forward stepwise selection
2 model_fds=regsubsets(y~.,data=question_3_data,nvmax=10,method='forward')
3 model_fds_summary=summary(model_fds)
```

```
1 model_fds_summary
```

```
Subset selection object
Call: regsubsets.formula(y ~ ., data = question_3_data, nvmax = 10,
    method = "forward")
10 Variables (and intercept)
  Forced in  Forced out
x      FALSE      FALSE
x.2     FALSE      FALSE
x.3     FALSE      FALSE
x.4     FALSE      FALSE
x.5     FALSE      FALSE
x.6     FALSE      FALSE
x.7     FALSE      FALSE
x.8     FALSE      FALSE
x.9     FALSE      FALSE
x.10    FALSE      FALSE
1 subsets of each size up to 10
Selection Algorithm: forward
      x   x.2 x.3 x.4 x.5 x.6 x.7 x.8 x.9 x.10
1 ( 1 )   *   *   *   *   *
2 ( 1 )   *   *   *   *   *
3 ( 1 )   *   *   *   *   *
4 ( 1 )   *   *   *   *   *
5 ( 1 )   *   *   *   *   *
6 ( 1 )   *   *   *   *   *
7 ( 1 )   *   *   *   *   *
8 ( 1 )   *   *   *   *   *
9 ( 1 )   *   *   *   *   *
10 ( 1 )  *   *   *   *   *
```

```
1 par(mfrow=c(3,1))
2 plot(model_fds_summary$cp,xlab='number of variables',ylab='cp',type='l')
3 plot(model_fds_summary$bic,xlab='number of variables',ylab='bic',type='l')
4 plot(model_fds_summary$adjr2,xlab='number of variables',ylab='adjr2',type='l')
```



```

1 indices<-c(which.min(model_fds_summary$cp),which.min(model_fds_summary$bic),which.max(model_fds_summary$adjr2))
2
3 for (index in indices) {
4   cat("Number of variables:", index, "\n")
5   cat("Coefficients:\n")
6   print(coef(model_fds, id = index))
7   cat("\n")
8 }

```

Number of variables: 3
Coefficients:
(Intercept) x x.2 x.3
0.9703939 1.9204462 2.9084569 4.0204363

Number of variables: 3
Coefficients:
(Intercept) x x.2 x.3
0.9703939 1.9204462 2.9084569 4.0204363

Number of variables: 4
Coefficients:
(Intercept) x x.2 x.3 x.6
1.040475559 1.928367169 2.746130084 4.023397035 0.009248115

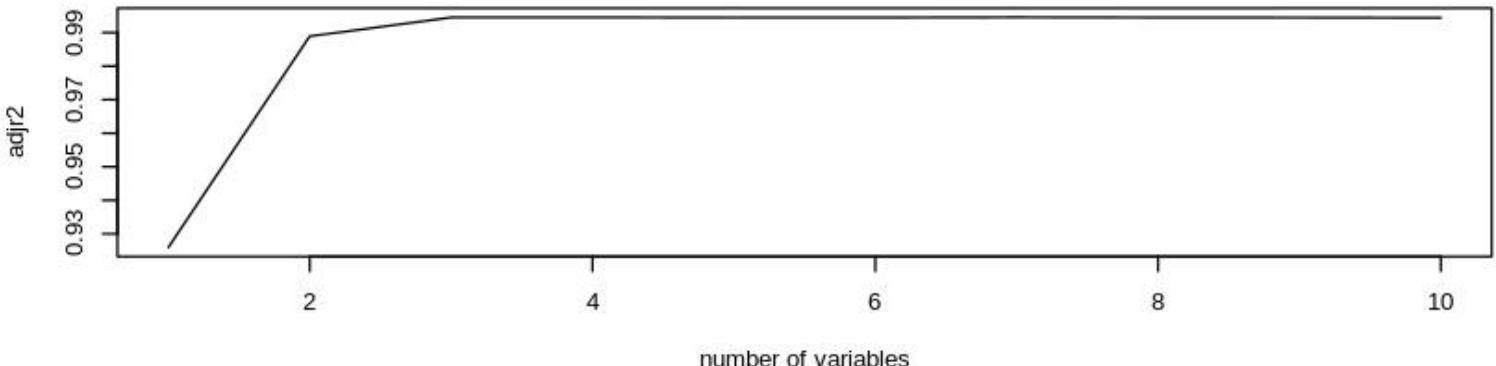
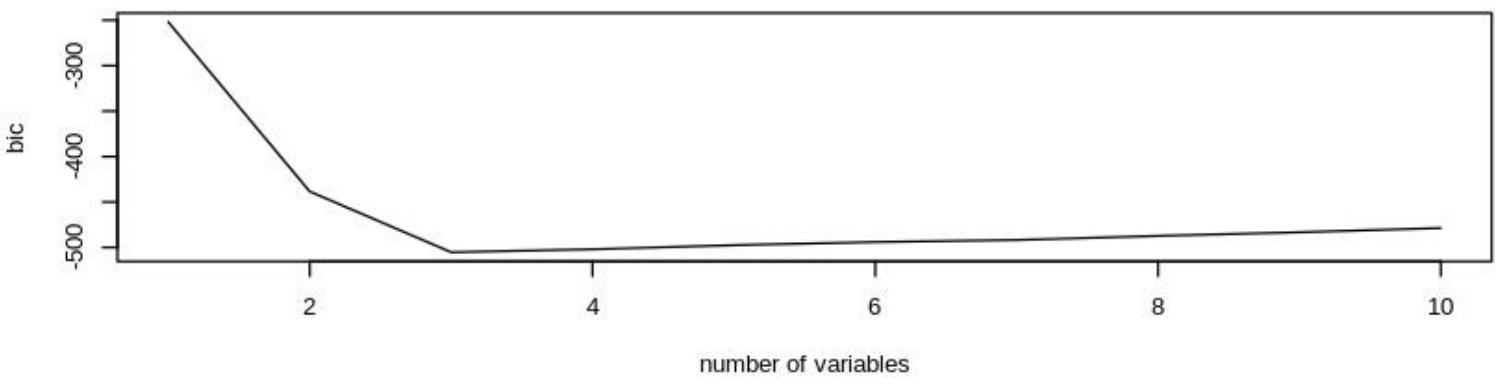
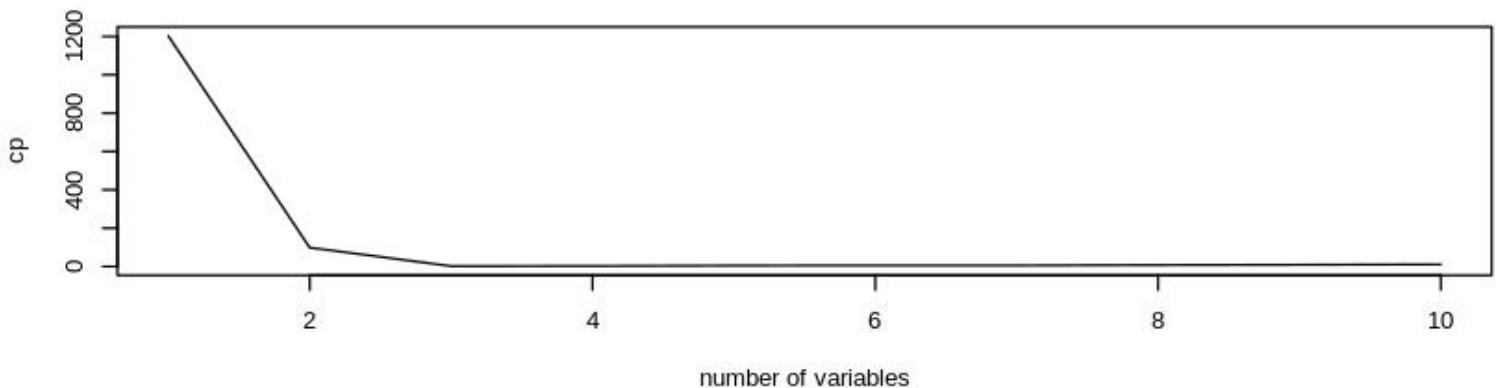
Backward Stepwise Selection

```
1 # backward stepwise selection
2 model_bds=regsubsets(y~.,data=question_3_data,nvmax=10,method='backward')
3 model_bds_summary=summary(model_bds)

] 1 model_bds_summary

Subset selection object
Call: regsubsets.formula(y ~ ., data = question_3_data, nvmax = 10,
    method = "backward")
10 Variables (and intercept)
  Forced in  Forced out
x      FALSE      FALSE
x.2    FALSE      FALSE
x.3    FALSE      FALSE
x.4    FALSE      FALSE
x.5    FALSE      FALSE
x.6    FALSE      FALSE
x.7    FALSE      FALSE
x.8    FALSE      FALSE
x.9    FALSE      FALSE
x.10   FALSE      FALSE
1 subsets of each size up to 10
Selection Algorithm: backward
      x  x.2 x.3 x.4 x.5 x.6 x.7 x.8 x.9 x.10
1 ( 1 )   *   *   *   *   *
2 ( 1 )   *   *   *   *   *
3 ( 1 )   *   *   *   *   *
4 ( 1 )   *   *   *   *   *
5 ( 1 )   *   *   *   *   *
6 ( 1 )   *   *   *   *   *
7 ( 1 )   *   *   *   *   *
8 ( 1 )   *   *   *   *   *
9 ( 1 )   *   *   *   *   *
10 ( 1 )  *   *   *   *   *

par(mfrow=c(3,1))
plot(model_bds_summary$cp,xlab='number of variables',ylab='cp',type='l')
plot(model_bds_summary$bic,xlab='number of variables',ylab='bic',type='l')
plot(model_bds_summary$adjr2,xlab='number of variables',ylab='adjr2',type='l')
```



```

1 indices<-c(which.min(model_bds_summary$cp),which.min(model_bds_summary$bic),which.max(model_bds_summary$adjr2))
2
3 for (index in indices) {
4   cat("Number of variables:",index,"\n")
5   cat("Coefficients:\n")
6   print(coef(model_bds,id=index))
7   cat("\n")
8 }
```

Number of variables: 3
Coefficients:
(Intercept) x x.2 x.3
0.9703939 1.9204462 2.9084569 4.0204363

Number of variables: 3
Coefficients:
(Intercept) x x.2 x.3
0.9703939 1.9204462 2.9084569 4.0204363

Number of variables: 7
Coefficients:
(Intercept) x x.2 x.3 x.4 x.6
0.7931750 1.8907484 5.2758848 4.0530514 -4.0243410 2.2138377
x.8 x.10
-0.4870719 0.0373538

All the models (Best subset, Forward Stepwise & Backward Stepwise selection method) gave the lowest Cp & BIC score for the model containing the predictors X , X^2 , and X^3 and only the Best subset and Backward Stepwise selection methods gave the highest value for adjusted R^2 for the model containing the predictors X , X^2 , X^3 , X^4 , X^6 , X^8 , X^{10} i.e model with 7 out of 10 predictors whereas Forward Stepwise selection method chose the model containing the predictors X , X^2 , X^3 , X^6 i.e model with only 4 out of 10 predictors.

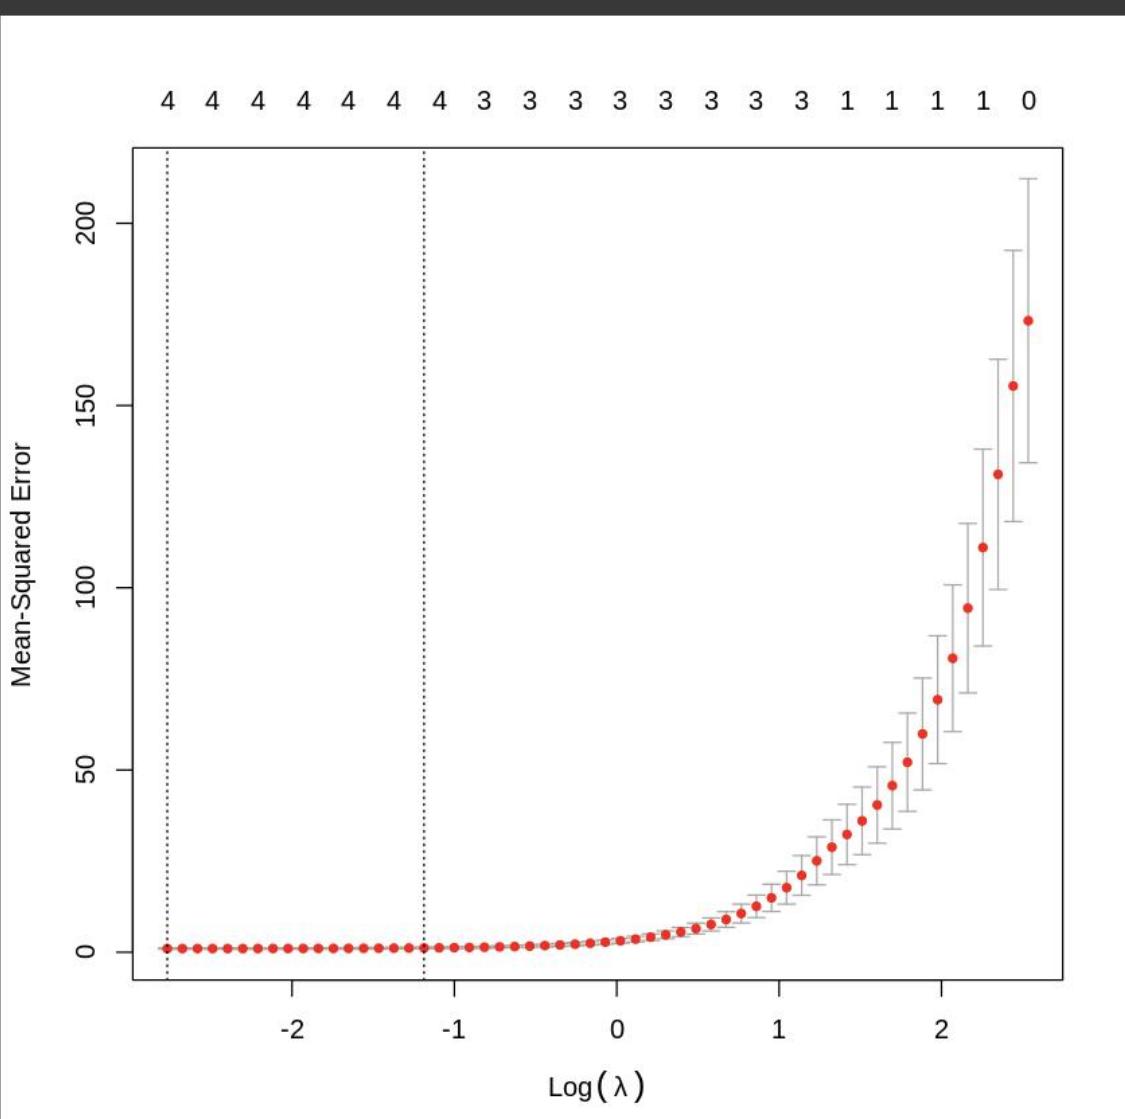
(e) Now fit a lasso model to the simulated data, again using X, X^2, \dots, X^{10} as predictors. Use cross-validation to select the optimal value of λ . Create plots of the cross-validation error as a function of λ . Report the resulting coefficient estimates, and discuss the results obtained.

```
library(glmnet)
```

```
xmat=model.matrix(y~.,data=question_3_data)[,-1]
model_lasso=cv.glmnet(xmat,y,alpha=1)
model_lasso_best_lambda=model_lasso$lambda.min
print(model_lasso_best_lambda)
```

0.06276895

```
1 plot(model_lasso)
```



```
1 model_best_lasso=glmnet(xmat,y,alpha=1)
2 predict(model_best_lasso,s=model_lasso_best_lambda,type="coefficients")

11 x 1 sparse Matrix of class "dgCMatrix"
      s1
(Intercept) 1.07946949
x           1.89142726
x.2         2.66822407
x.3         4.01088941
x.4         0.04840102
x.5         .
x.6         .
x.7         .
x.8         .
x.9         .
x.10        .
```

After performing the cross-validation using the lasso method we find that lasso puts more weight on X, X^2 , X^3 , and X^4 predictors stating that other predictors don't contribute much to the prediction of response Y.

(f) Now generate a response vector Y according to the model $Y = \beta_0 + \beta_7 X^7 + \varepsilon$, and perform best subset selection and the lasso. Discuss the results obtained.

beta_7=6

`y = beta_0 + (beta_7 * (x^7)) + e_noise`

```
question_3_f_data=data.frame(x,x^2,x^3,x^4,x^5,x^6,x^7,x^8,x^9,x^10,y)
```

```
1 # best subset selection method
2 model_bss_q3f=regsubsets(y~.,data=question_3_f_data,nvmax=10)
3
4 model_bss_q3f_summary=summary(model_bss_q3f)
5 model_bss_q3f_summary
```

Subset selection object

```
Call: regsubsets.formula(y ~ ., data = question_3_f_data, nvmax = 10)
10 Variables (and intercept)
```

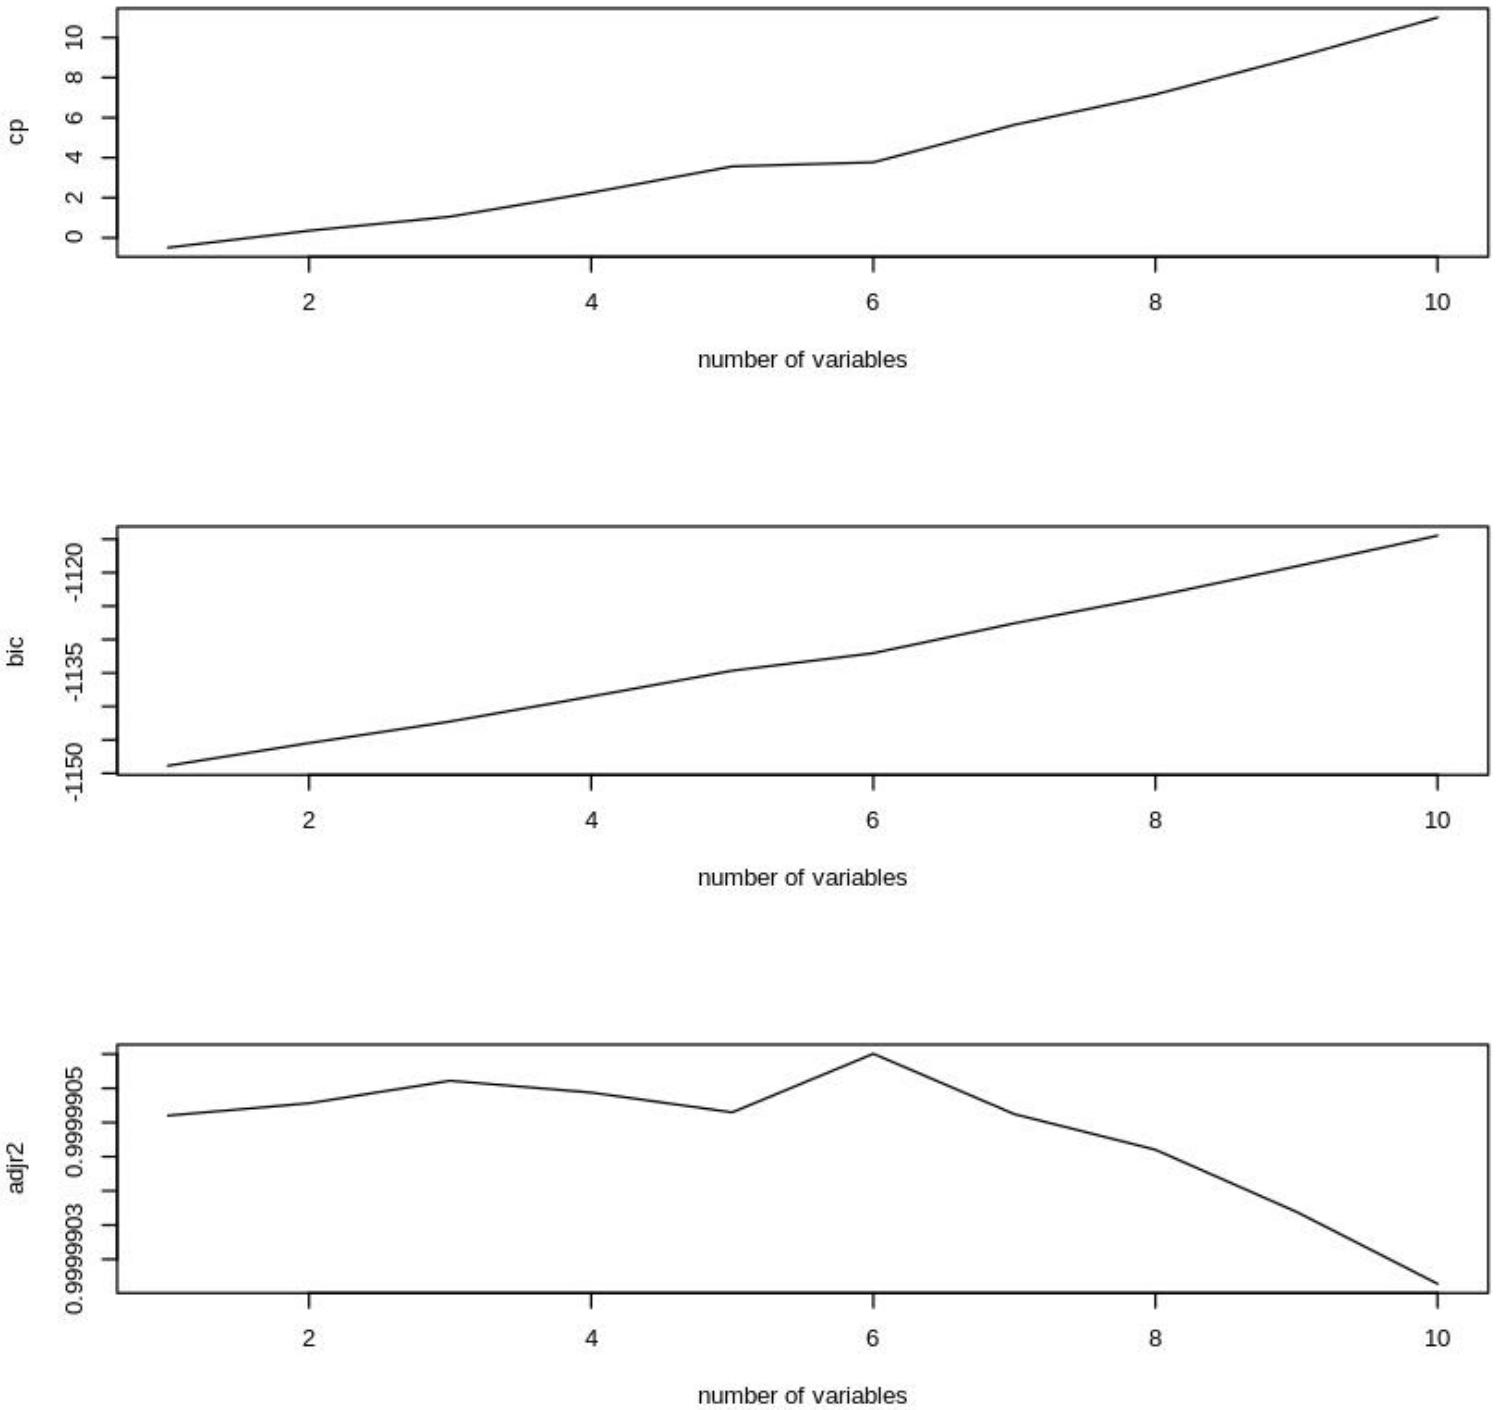
Forced in Forced out

x	FALSE	FALSE
x.2	FALSE	FALSE
x.3	FALSE	FALSE
x.4	FALSE	FALSE
x.5	FALSE	FALSE
x.6	FALSE	FALSE
x.7	FALSE	FALSE
x.8	FALSE	FALSE
x.9	FALSE	FALSE
x.10	FALSE	FALSE

1 subsets of each size up to 10

Selection Algorithm: exhaustive

```
par(mfrow=c(3,1))
plot(model_bss_q3f_summary$cp,xlab='number of variables',ylab='cp',type='l')
plot(model_bss_q3f_summary$bic,xlab='number of variables',ylab='bic',type='l')
plot(model_bss_q3f_summary$adjr2,xlab='number of variables',ylab='adjr2',type='l')
```



```

1 indices<-c(which.min(model_bss_q3f_summary$cp),which.min(model_bss_q3f_summary$bic),which.max(model_bss_q3f_summary$adjr2))
2
3 for (index in indices) {
4   cat("Number of variables:",index,"\n")
5   cat("Coefficients:\n")
6   print(coef(model_bss_q3f,id=index))
7   cat("\n")
8 }

Number of variables: 1
Coefficients:
(Intercept)      x.7
 0.8932512  5.9997045

Number of variables: 1
Coefficients:
(Intercept)      x.7
 0.8932512  5.9997045

Number of variables: 6
Coefficients:
(Intercept)      x.2      x.4      x.6      x.7      x.8
 0.79358642  2.30456710 -4.10732161  2.27693175  6.00151236 -0.50477246
           x.10
 0.03899314

```

```

# lasso selection method
xmat=model.matrix(y~.,data=question_3_f_data)[,-1]
model_lasso=cv.glmnet(xmat,y,alpha=1)
model_lasso_best_lambda=model_lasso$lambda.min
print(model_lasso_best_lambda)

```

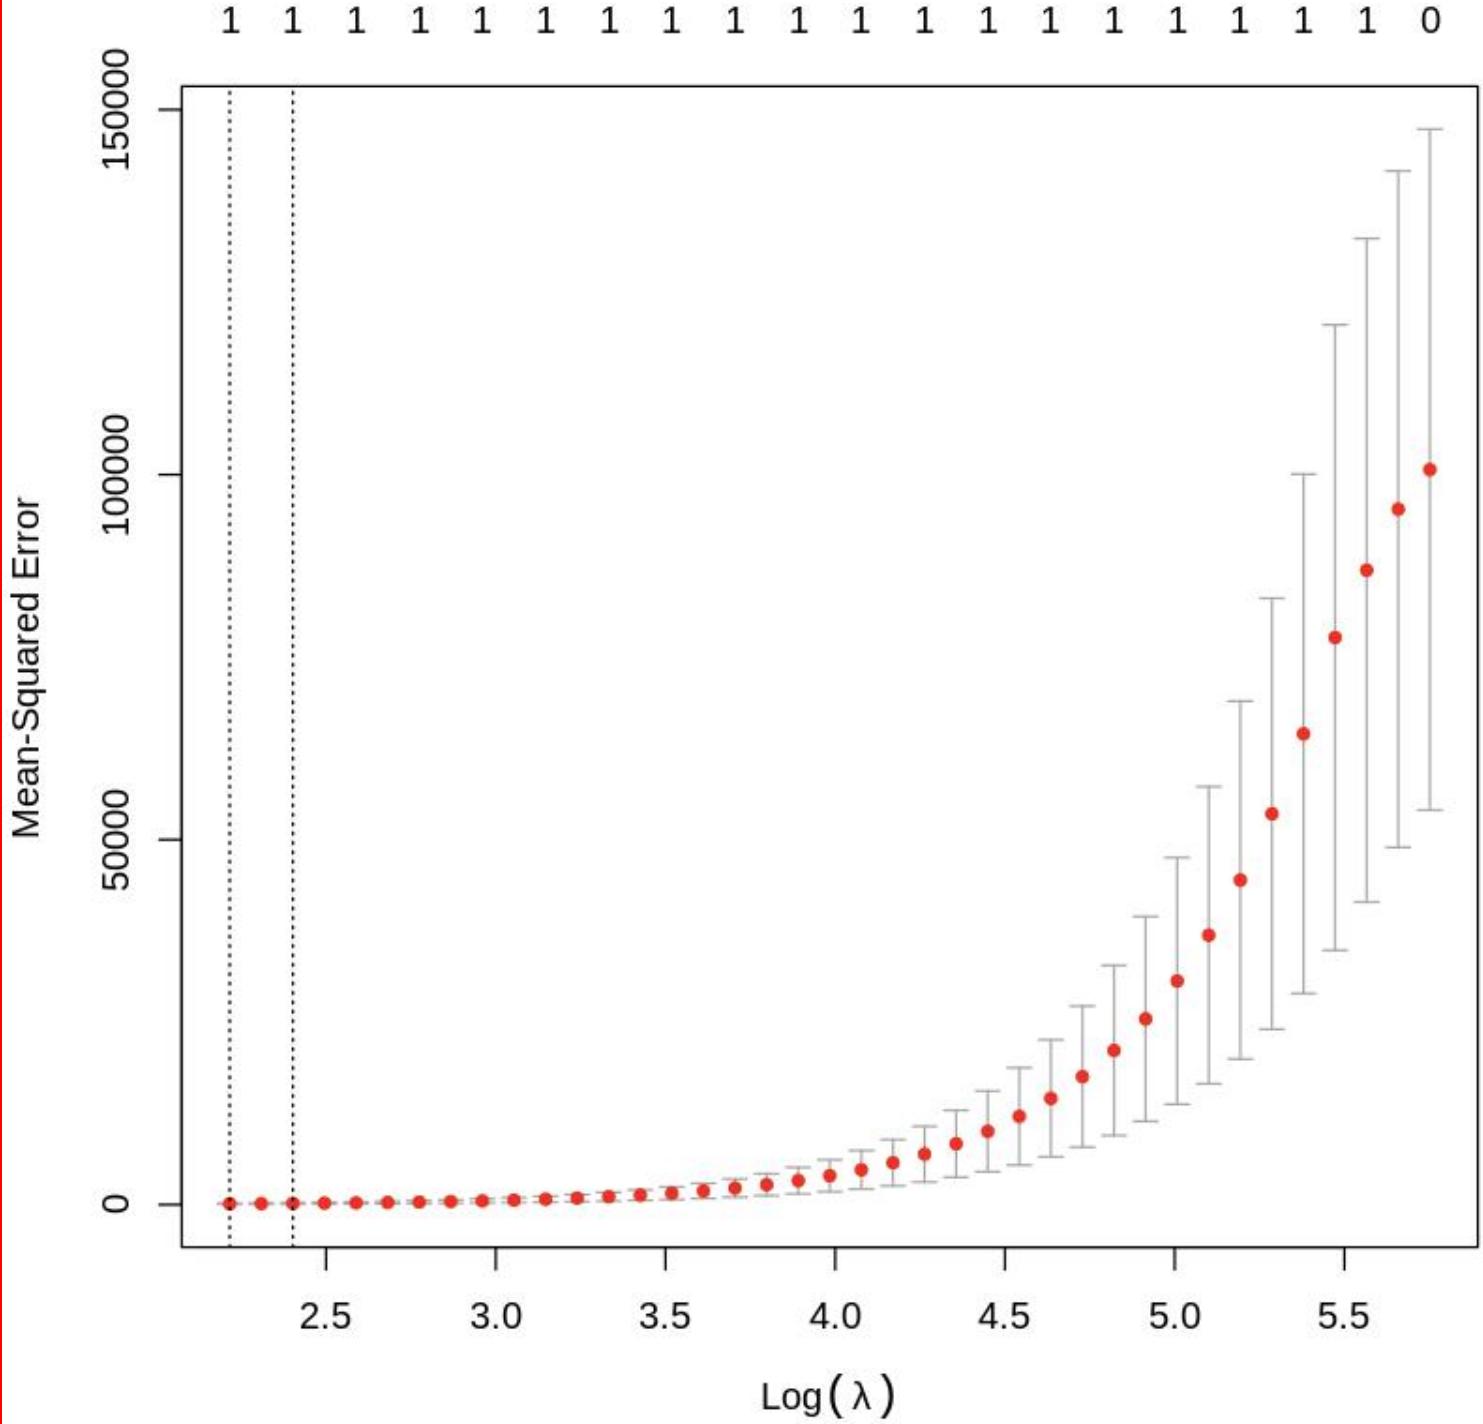
```

1 model_best_lasso=glmnet(xmat,y,alpha=1)
2 predict(model_best_lasso,s=model_lasso_best_lambda,type="coefficients")

11 x 1 sparse Matrix of class "dgCMatrix"
           s1
(Intercept) 1.365508
x            .
x.2          .
x.3          .
x.4          .
x.5          .
x.6          .
x.7          5.824810
x.8          .
x.9          .
x.10         .

```

```
1 plot(model_lasso)
```



The Best subset selection method found that for Cp and BIC score it was minimum for predictor X^3 and for maximum adjusted R^2 the best model was the one with 6 predictors namely X , X^2 , X^3 , X^5 , X^6 and X^9 . However, after performing the cross-validation using the lasso method, we find that lasso puts more weight only on the X^7 predictor stating that other predictors don't contribute much to the prediction of response Y.

4. In this exercise, we will predict the number of applications received using the other variables in the College data set.

```
library(ISLR)

set.seed(123)

sum(is.na(College))
```

(a) Split the data set into a training set and a test set.

```
1 round(nrow(College)*0.2),0)
155

1 test_index=sample(nrow(College),round((nrow(College)*0.2),0),replace=FALSE)
2 train_set=College[-test_index,]
3 test_set=College[test_index,]

1 c(nrow(train_set),nrow(test_set),nrow(train_set)+nrow(test_set),nrow(College))
622 · 155 · 777 · 777

1 names(College)
'Private' · 'Apps' · 'Accept' · 'Enroll' · 'Top10perc' · 'Top25perc' · 'F.Undergrad' · 'P.Undergrad' · 'Outstate' · 'Room.Board' · 'Books' · 'Personal' · 'PhD' · 'Terminal' · 'S.F.Ratio' · 'perc.alumni' · 'Expend' · 'Grad.Rate'
```

(b) Fit a linear model using least squares on the training set, and report the test error obtained.

```
1 lm_model=lm(Apps~.,data=train_set)
2 lm_model_pred=predict(lm_model,test_set)
3 # RSS
4 print(mean((test_set[, "Apps"]-lm_model_pred)^2))

[1] 726062.2
```

(c) Fit a ridge regression model on the training set, with λ chosen by cross-validation. Report the test error obtained.

```
library(glmnet)
train_xmat=model.matrix(Apps~.,data=train_set)
test_xmat=model.matrix(Apps~.,data=test_set)

1 # Ridge
2 model_ridge=cv.glmnet(train_xmat,train_set[, 'Apps'],alpha=0)
3 model_ridge_best_lambda=model_ridge$lambda.min
4 print(model_ridge_best_lambda)

[1] 387.5921

ridge_model_pred=predict(model_ridge,test_xmat,s=model_ridge_best_lambda)
# RSS for Ridge
print(mean((test_set[, "Apps"]-ridge_model_pred)^2))

799031.9

1 plot(model_ridge)
```

The figure is a line graph titled 'plot(model_ridge)'. The x-axis is labeled 'Log(λ)' and has tick marks at 6, 8, 10, 12, 14, and 17. The y-axis is labeled 'Mean-Squared Error' and has tick marks at 5.0e+06, 1.0e+07, 1.5e+07, and 2.0e+07. A red line represents the mean-squared error for different values of λ . The error starts at approximately 1.5e+06 when $\log(\lambda) = 6$, remains relatively flat until $\log(\lambda) \approx 7$, then rises sharply, reaching about 1.8e+07 at $\log(\lambda) = 17$. A vertical dashed line is drawn at $\log(\lambda) \approx 7.2$, corresponding to the minimum point of the red curve.

(d) Fit a lasso model on the training set, with λ chosen by cross-validation. Report the test error obtained, along with the number of non-zero coefficient estimates.

Lasso

```
model_lasso=cv.glmnet(train_xmat,train_set[, 'Apps'],alpha=1)
model_lasso_best_lambda=model_lasso$lambda.min
print(model_lasso_best_lambda)
```

12.11503

```
lasso_model_pred=predict(model_lasso,test_xmat,s=model_lasso_best_lambda)
# RSS for Lasso
print(mean((test_set[, "Apps"]-lasso_model_pred)^2))
```

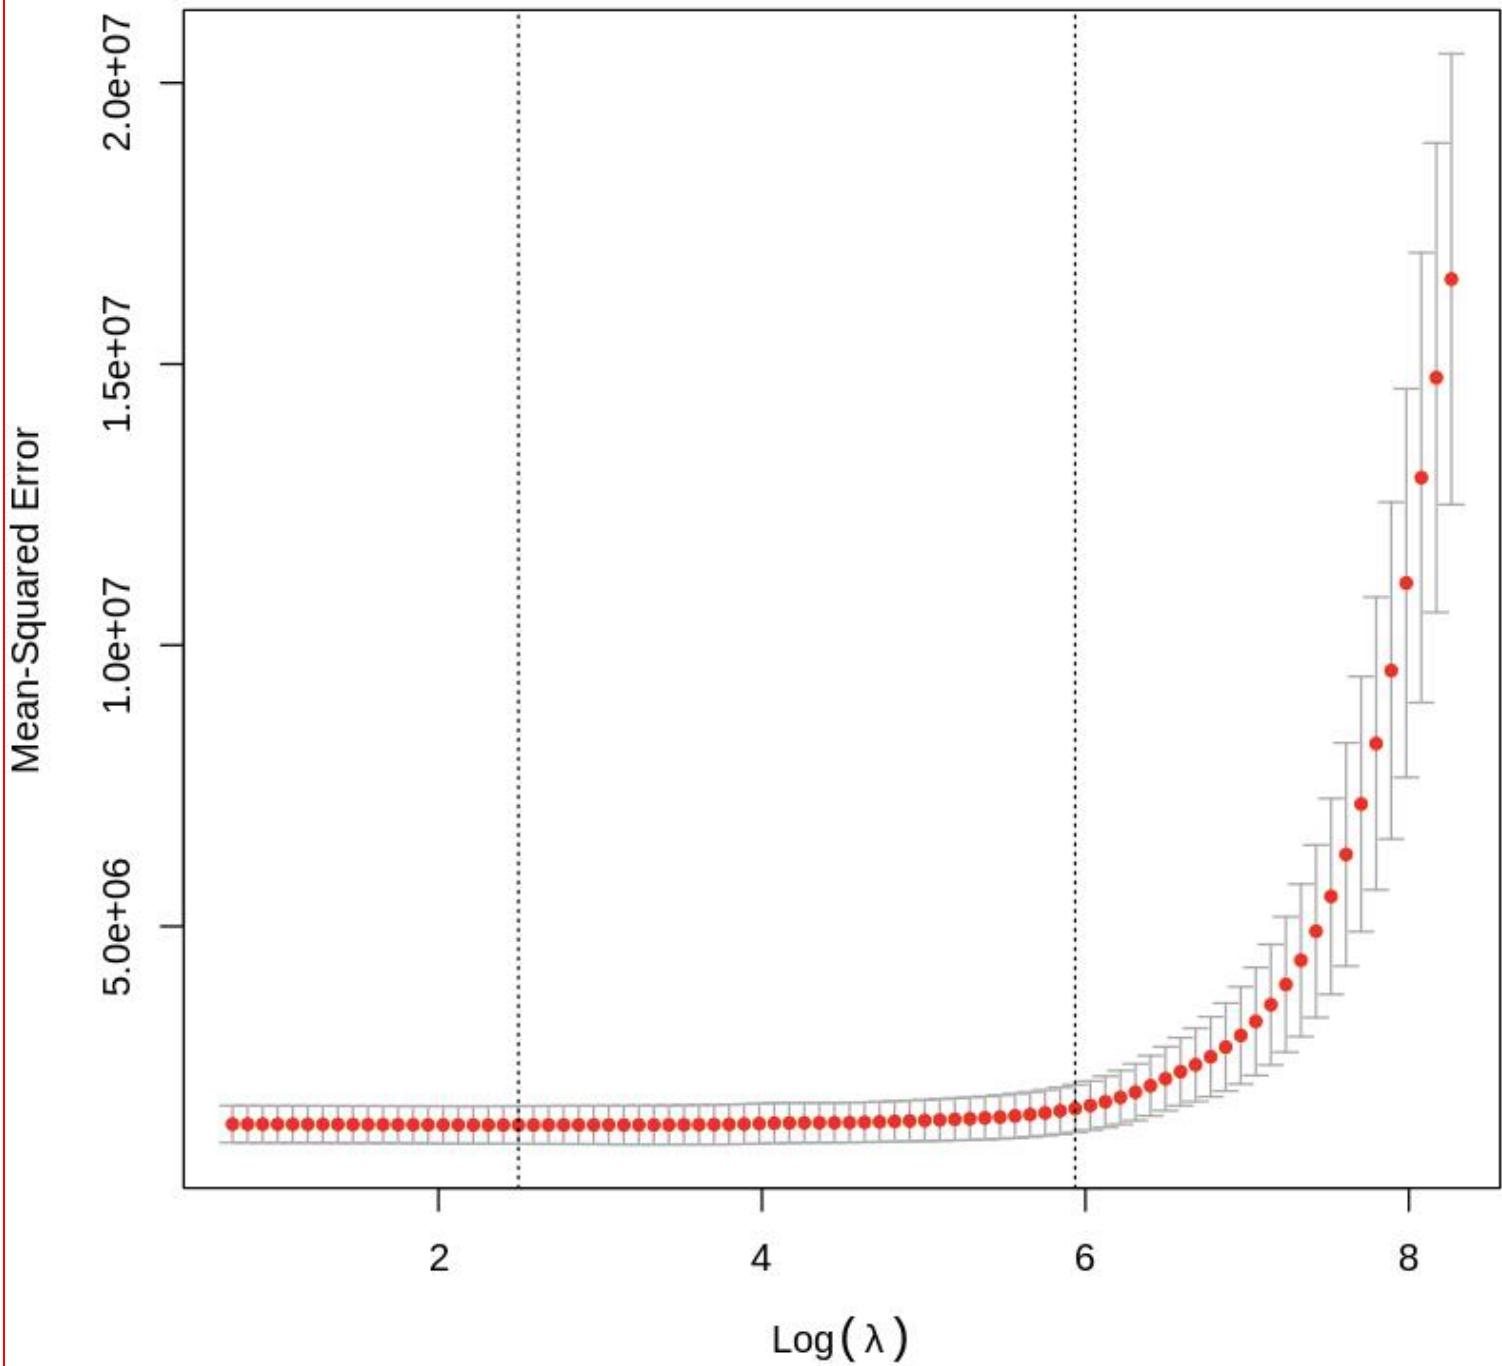
750801.5

```
1 predict(model_lasso,test_xmat,s=model_lasso_best_lambda,type='coefficients')

19 x 1 sparse Matrix of class "dgCMatrix"
      s1
(Intercept) -6.183583e+02
(Intercept) .
PrivateYes   -4.079245e+02
Accept        1.527555e+00
Enroll       -4.085465e-01
Top10perc    3.813867e+01
Top25perc    -5.484379e+00
F.Undergrad   .
P.Undergrad   3.469064e-02
Outstate     -7.896075e-02
Room.Board    1.138274e-01
Books         1.426352e-01
Personal      2.533014e-03
PhD          -6.384711e+00
Terminal      -2.743584e+00
S.F.Ratio     5.759028e+00
perc.alumni   .
Expend        7.859546e-02
Grad.Rate     7.305107e+00
```

```
1 plot(model_lasso)
```

```
17 17 16 15 15 14 11 9 4 4 3 3 3 2 1 1 1 1 1 1
```



(e) Fit a PCR model on the training set, with M chosen by cross-validation. Report the test error obtained, along with the value of M selected by cross-validation.

```
1 library(pls)
```

Attaching package: 'pls'

The following object is masked from 'package:stats':

loadings

```
1 set.seed(123)
```

```
1 model_pcr=pcr(Apps~.,data=train_set,scale=TRUE,validation="CV")
```

```
1 summary(model_pcr)
```

Data: X dimension: 622 17
Y dimension: 622 1

Fit method: svdpc

Number of components considered: 17

VALIDATION: RMSEP

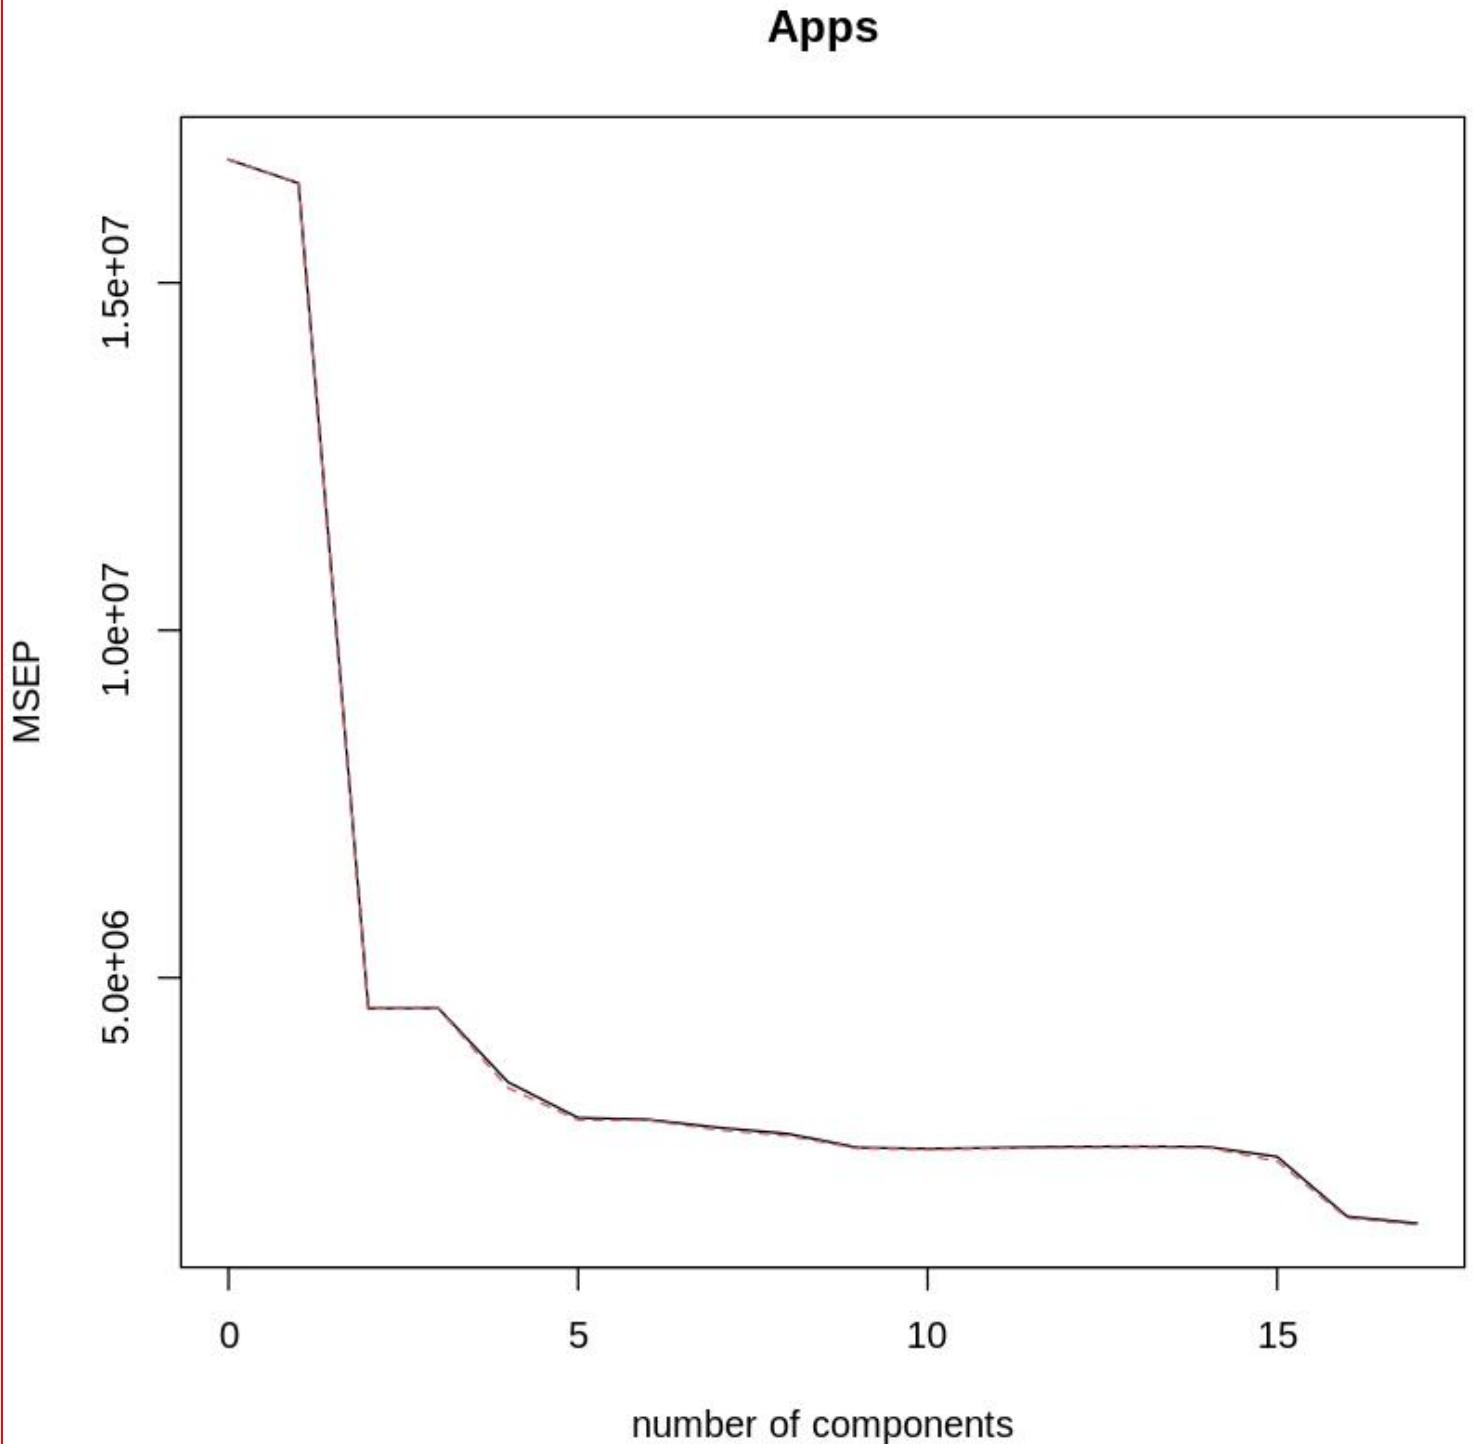
Cross-validated using 10 random segments.

	(Intercept)	1 comps	2 comps	3 comps	4 comps	5 comps	6 comps
CV	4096	4054	2135	2136	1870	1728	1720
adjCV	4096	4054	2133	2134	1847	1719	1717
	7 comps	8 comps	9 comps	10 comps	11 comps	12 comps	13 comps
CV	1686	1660	1599	1594	1598	1603	1604
adjCV	1677	1651	1595	1591	1595	1599	1601
	14 comps	15 comps	16 comps	17 comps			
CV	1603	1557	1251	1211			
adjCV	1599	1536	1242	1202			

TRAINING: % variance explained

	1 comps	2 comps	3 comps	4 comps	5 comps	6 comps	7 comps	8 comps
X	30.699	56.57	63.82	69.58	74.97	79.92	83.57	87.14
Apps	2.599	73.20	73.38	80.31	83.28	83.35	84.43	84.84
	9 comps	10 comps	11 comps	12 comps	13 comps	14 comps	15 comps	
X	90.29	92.80	94.90	96.74	97.85	98.73	99.36	
Apps	85.60	85.78	85.78	85.80	85.81	85.83	90.92	
	16 comps	17 comps						
X	99.84	100.00						
Apps	92.80	93.15						

```
1 validationplot(model_pcr, val.type="MSEP")
```



```
1 names(model_pcr)
```

```
'coefficients' · 'scores' · 'loadings' · 'Yloadings' · 'projection' · 'Xmeans' · 'Ymeans' · 'fitted.values' · 'residuals' · 'Xvar' · 'Xtotvar' · 'fit.time' · 'ncomp' · 'method' · 'center' · 'scale' · 'validation' · 'call' · 'terms' · 'model'
```

```

1 model_pcr$validation
3529.23722884627 · 7972.00684733083 · 5039.12191363883 · 3308.79482709097 · 16533.1812358161 · 1709.13598804866 · 1007.40992271018 · 542.128685230945 · 1444.10759381807 · 730.509112834031 · 1383.05215739374 ·
1686.21727223182 · 3438.82062735123 · 1253.86519252245 · 8883.16886052385 · 12015.3267268696 · 1926.0371268543 · 1047.44884471543 · -72.8654593176093 · 325.828909125968 · 4507.19263323483 · 557.642537257667 ·
3035.27391799075 · 902.504881038608 · 4115.38347015494 · 4484.77744319732 · 1962.9891202398 · 1655.97537484267 · 3338.42600729071 · 1266.53044352338 · 10513.0239755484 · 1744.66081581826 · 1993.73959707934 ·
5166.14105529183 · 2706.9813173143 · 471.218508313124 · 1139.50984368088 · 615.292113451915 · 1723.56342380075 · 2321.53803655795 · 2212.04538431593 · 637.56085540921 · 1496.29465765044 · 1674.24857948719 ·
1704.87417025867 · 2174.30812133138 · 499.661515128412 · 345.41194326638 · 3974.31701344962 · -103.826497963051 · 1058.7857493162 · 2692.15176175419 · 19.7900187031174 · 2531.17673590743 · 1888.20181517252 ·
3706.19026428798 · 2671.341699332398 · 3183.19725119152 · 2919.75077479317

$coefficients
NULL
$gammas
NULL
$PRESSO
  Apps: 10433066442.2424
$PRESS

A matrix: 1 × 17 of type dbl
  1 comps   2 comps   3 comps   4 comps   5 comps   6 comps   7 comps   8 comps   9 comps   10 comps   11 comps   12 comps   13 comps   14 comps   15 comps   16 comps
 Apps 10221327994 2836331306 2837640890 2174425480 1856403269 1841060167 1768814640 1713568318 1590008381 1580372354 1588545069 1597583639 1600240595 1597348814 1508141487 973559455

$adj
A matrix: 1 × 17 of type dbl
  1 comps   2 comps   3 comps   4 comps   5 comps   6 comps   7 comps   8 comps   9 comps   10 comps   11 comps   12 comps   13 comps   14 comps   15 comps   16 comps   17 comps
 Apps 16279889 4489677 4457159 3374094 2825616 2795139 2635904 2564797 2420648 2386627 2387357 2385472 2382463 2379599 1583036 1227043 1165422

$segments
$V1
471 · 48 · 481 · 8 · 531 · 565 · 376 · 570 · 407 · 321 · 151 · 410 · 350 · 600 · 604 · 144 · 28 · 476 · 422 · 555 · 27 · 461 · 506 · 78 · 115 · 164 · 282 · 230 · 91 · 23 · 548 · 427 · 236 · 606 · 402 · 446 · 95 · 12 · 349 · 553 · 360 · 136 · 47 · 179 · 362 ·
283 · 263 · 429 · 547 · 605 · 88 · 273 · 249 · 279 · 552 · 334 · 431 · 259 · 260 · 59 · 254 · 81 · 219
$V2
227 · 482 · 250 · 258 · 488 · 329 · 84 · 200 · 595 · 9 · 622 · 472 · 314 · 142 · 573 · 308 · 157 · 439 · 152 · 480 · 201 · 327 · 504 · 158 · 238 · 382 · 292 · 223 · 89 · 216 · 435 · 515 · 24 · 239 · 45 · 202 · 498 · 128 · 354 · 507 · 1 · 585 · 549 · 578 · 15 ·
54 · 13 · 558 · 590 · 160 · 588 · 325 · 31 · 317 · 306 · 309 · 141 · 405 · 120 · 340 · 228 · 196 · 131
$V3
315 · 248 · 597 · 262 · 271 · 339 · 617 · 237 · 391 · 281 · 344 · 572 · 217 · 229 · 86 · 559 · 300 · 372 · 208 · 513 · 211 · 366 · 364 · 156 · 612 · 567 · 448 · 101 · 143 · 277 · 154 · 181 · 270 · 557 · 520 · 37 · 244 · 176 · 501 · 320 · 601 · 621 · 71 ·
524 · 177 · 73 · 414 · 603 · 189 · 577 · 359 · 301 · 492 · 430 · 502 · 433 · 452 · 68 · 14 · 499 · 442 · 82
$V4
243 · 105 · 330 · 529 · 398 · 583 · 554 · 591 · 290 · 40 · 168 · 224 · 184 · 375 · 434 · 35 · 310 · 546 · 75 · 100 · 475 · 34 · 215 · 399 · 138 · 335 · 313 · 145 · 123 · 275 · 419 · 16 · 540 · 522 · 367 · 495 · 449 · 358 · 102 · 337 · 512 · 294 · 284 · 107 ·
97 · 22 · 134 · 550 · 274 · 246 · 118 · 178 · 21 · 85 · 66 · 108 · 261 · 374 · 121 · 392 · 267 · 497
$V5
265 · 171 · 195 · 206 · 218 · 247 · 194 · 380 · 287 · 148 · 161 · 592 · 608 · 64 · 17 · 438 · 153 · 496 · 87 · 425 · 418 · 381 · 390 · 186 · 462 · 474 · 3 · 133 · 413 · 420 · 569 · 403 · 109 · 467 · 83 · 511 · 491 · 288 · 613 · 543 · 443 · 150 · 336 · 357 ·
338 · 20 · 453 · 571 · 599 · 233 · 33 · 464 · 62 · 445 · 296 · 289 · 38 · 245 · 389 · 345 · 485 · 57
$V6
535 · 459 · 99 · 199 · 155 · 370 · 460 · 616 · 466 · 18 · 421 · 400 · 451 · 242 · 137 · 26 · 285 · 556 · 94 · 25 · 191 · 517 · 103 · 286 · 204 · 174 · 386 · 371 · 30 · 328 · 441 · 240 · 190 · 46 · 477 · 6 · 4 · 377 · 394 · 395 · 52 · 415 · 117 · 165 · 519 ·
114 · 92 · 416 · 307 · 505 · 125 · 50 · 463 · 104 · 135 · 257 · 276 · 331 · 113 · 412 · 368 · 106
$V7
112 · 516 · 98 · 607 · 209 · 530 · 77 · 5 · 566 · 579 · 444 · 39 · 56 · 280 · 163 · 269 · 483 · 593 · 304 · 437 · 343 · 303 · 80 · 500 · 494 · 426 · 11 · 369 · 393 · 197 · 564 · 582 · 116 · 332 · 298 · 252 · 319 · 188 · 611 · 455 · 7 · 436 · 147 · 423 · 60 ·
542 · 450 · 322 · 586 · 212 · 609 · 251 · 378 · 373 · 574 · 232 · 132 · 509 · 563 · 299 · 231 · 544
```

```

1 model_pcr_pred=predict(model_pcr,test_set,ncomp=17)
2 round((mean((test_set[, 'Apps'] - model_pcr_pred)^2)),0)
```

726062

(f) Fit a PLS model on the training set, with M chosen by cross-validation. Report the test error obtained, along with the value of M selected by cross-validation.

```
1 model_plsr=plsR(Apps~, data=train_set, scale=TRUE, validation="CV")
```

```
1 summary(model_plsr)
```

Data: X dimension: 622 17

Y dimension: 622 1

Fit method: kernelpls

Number of components considered: 17

VALIDATION: RMSEP

Cross-validated using 10 random segments.

	(Intercept)	1 comps	2 comps	3 comps	4 comps	5 comps	6 comps
--	-------------	---------	---------	---------	---------	---------	---------

CV	4096	1971	1626	1533	1431	1267	1208
----	------	------	------	------	------	------	------

adjCV	4096	1968	1622	1528	1416	1255	1201
-------	------	------	------	------	------	------	------

	7 comps	8 comps	9 comps	10 comps	11 comps	12 comps	13 comps
--	---------	---------	---------	----------	----------	----------	----------

CV	1205	1196	1195	1193	1190	1189	1188
----	------	------	------	------	------	------	------

adjCV	1198	1189	1187	1186	1184	1182	1181
-------	------	------	------	------	------	------	------

	14 comps	15 comps	16 comps	17 comps
--	----------	----------	----------	----------

CV	1187	1187	1187	1187
----	------	------	------	------

adjCV	1181	1180	1180	1180
-------	------	------	------	------

TRAINING: % variance explained

	1 comps	2 comps	3 comps	4 comps	5 comps	6 comps	7 comps	8 comps
--	---------	---------	---------	---------	---------	---------	---------	---------

X	25.96	38.18	61.83	64.80	68.33	73.36	77.46	80.72
---	-------	-------	-------	-------	-------	-------	-------	-------

Apps	78.02	85.59	87.61	90.65	92.49	92.93	92.98	93.03
------	-------	-------	-------	-------	-------	-------	-------	-------

	9 comps	10 comps	11 comps	12 comps	13 comps	14 comps	15 comps
--	---------	----------	----------	----------	----------	----------	----------

X	82.54	85.20	88.15	90.63	92.25	94.38	96.96
---	-------	-------	-------	-------	-------	-------	-------

Apps	93.10	93.12	93.13	93.14	93.15	93.15	93.15
------	-------	-------	-------	-------	-------	-------	-------

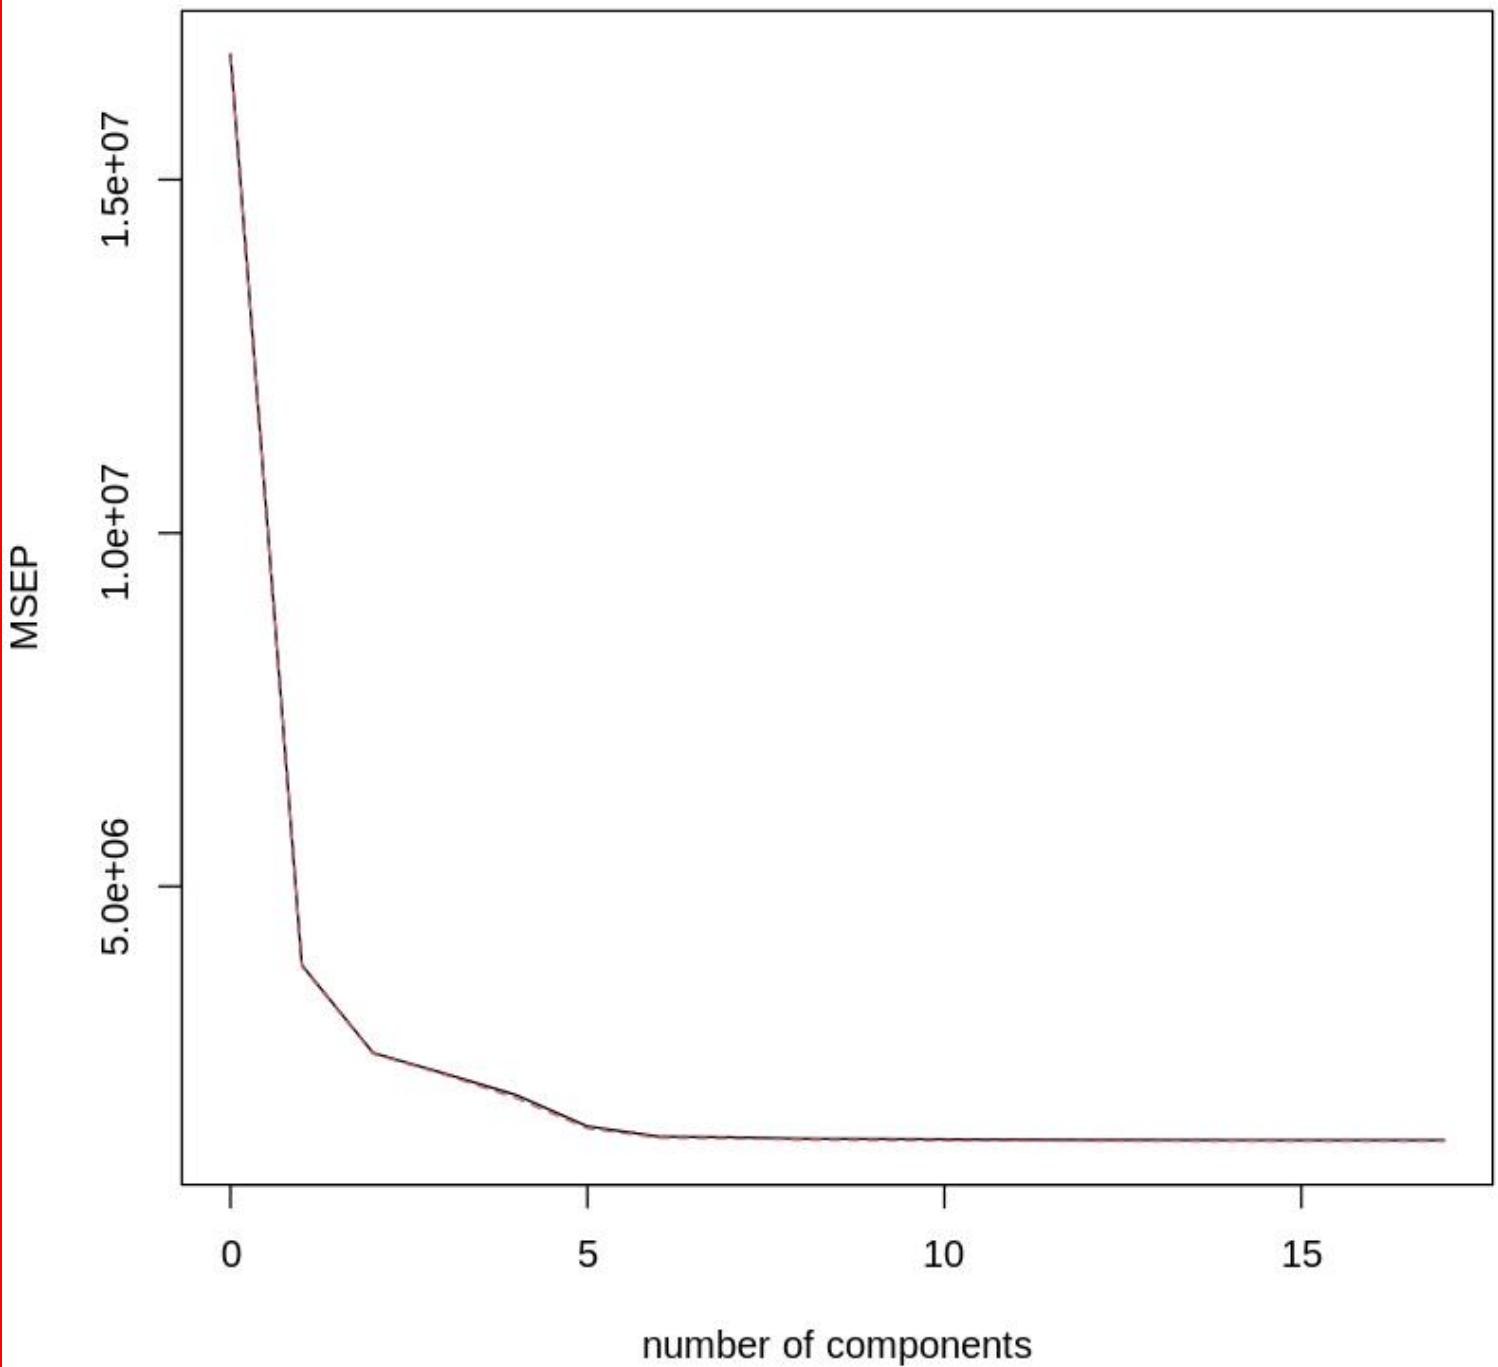
	16 comps	17 comps
--	----------	----------

X	97.95	100.00
---	-------	--------

Apps	93.15	93.15
------	-------	-------

```
1 validationplot(model_plsr, val.type="MSEP")
```

Apps



```

1 model_plsr$validation
2
3 $method
4   'CV'
5 $pred
6
7 2073.59008260573 · 57.2099396276953 · 88.9121827620693 · 2933.66928645074 · 650.294420285835 · 1423.65038223042 · 3060.97696631763 · 1534.30524991411 · -875.050469892874 · 2729.75427195453 · 2536.83701337515 ·
8 60.8953827705304 · 740.808211878373 · 4837.86469011346 · -21.5092947986836 · 66.6255892632748 · 1980.38818134788 · 6764.77162598306 · 13509.5261099642 · 1827.21486086216 · 2862.08058834378 · 10137.3606636982 ·
9 -60.2989470563339 · 2299.36400322152 · 571.836093674192 · 2296.09740843603 · 516.171691644397 · 2919.92975120124 · -510.443526366508 · 3666.23017559483 · 4384.4208011017 · 1751.17745739585 · 817.358595742962 ·
10 900.034220741616 · 1680.51344022965 · 1489.5300617588 · -297.830729782615 · -512.7529295433986 · 3743.29755333456 · 2426.21650468125 · -1182.15141627074 · -128.471422671984 · 2431.35935504034 · 39.0606159681229 ·
11 3880.01207391842 · -210.01145216567 · -429.527116218209 · 13057.5833142104 · 5651.335323156349 · -8697.24617778204 · 497.7001050536187 · 4915.90776757338 · 344.582457720395 · -616.553209112386 · 11581.5117416502 ·
12 7494.75517368663 · 5727.37513319961 · 3072.99452690043 · -98.4767222994201 · 1261.18540877082 · 890.769990063968 · 7697.32867356495 · 7111.58270848894 · 2252.81409712199 · 1717.32768081584 · -2214.74392880023 ·
13 4059.886651009452 · 6644.03777529298 · 5416.540795868104 · 2112.74666310087 · 776.90853842857 · 2984.13063093042 · -3.53144235937543 · 1050.31130653721 · 346.036944234849 · -450.887950527494 · 1793.1804341345 ·
14 13.936832357821 · 1848.45724460864 · 3730.34987188496 · -4971.0160783294 · 4198.31588107414 · -1175.26405585634 · 3085.78303705047 · 1689.22565838588 · 1764.20322518812 · 1410.84735151323 · 626.234543412661 ·
15 -91.8864000950375 · 3842.42292328646 · 9.160.53208956952 · -427.924376042134 · -213.34838119414 · 593.453007262012 · 96.95535862133516 · 3826.938609090453 · 1543.04423352658 · -576.617869221361 · 146.714758586201 · -1267.39686586329 ·
16 -944.28978753886 · 933.752886744391 · 2854.58075409339 · 1717.48205024486 · 233.4237052736361 · -1191.46962573278747 · 450.86742550805 · 6376.06253807705 · 2617.99882990733 · 9367.57363846405 ·
17 901.1045161686835 · 6160.53208956952 · -427.924376042134 · -213.34838119414 · 593.453007262012 · 96.95535862133516 · 3826.938609090453 · 1543.04423352658 · -576.617869221361 · 146.714758586201 · -1267.39686586329 ·
18 -944.28978753886 · 933.752886744391 · 2854.58075409339 · 1717.48205024486 · 233.4237052736361 · -1191.46962573278747 · 450.86742550805 · 6376.06253807705 · 2617.99882990733 · 9367.57363846405 ·
19 7962.10047609445 · 2356.59604477856 · 243.942605759206 · 2640.97582787497 · 4520.32787645597 · -92.0683386913452 · -259.830525601647 · 2365.90097124608 · 1834.38828302682 · 2039.07751922954 · 634.304487034391 ·
20 1909.57324654616 · 1643.96892440449 · 1620.11177474743 · 7740.81524820692 · -3776.05780728977 · 1354.60972518165 · 4551.36415413253 · 709.0960816299 · -498.148860669785 · 121.505921509973 · 3320.72480832 ·
21 8851.91281376503 · 1088.80442523646 · 12178.4427681081 · -775.512136415984 · 4845.14822524977 · 4286.24275060326 · 1046.723476170708 · 1045.46561621815 · 1084.8768522354 · 4344.21759202625 · 1083.32855866281 ·
22 -275.19434231767 · -303.200444326031 · 7261.77085395001 · 6678.63570754221 · 104.86149035161 · 7228.73038438624 · -372.819236030287 · 3700.783.20616759 · -1653.49589015081 · 3006.83135367814 · -253.907919714269 ·
23 2360.49580737221 · -612.326623609297 · -152.294240949243 · -210.5331955723 · -210.82610477394 · 2210.826120477394 · -607.626120477394 · 150.5331955723 · -152.294240949243 ·
24 3542.7182223324 · 1546.35786964352 · 753.919622891721 · -2327.06646882323 · -8320.06179014564 · 5197.615148418899 · 90.090312085526 · -2597.91442074948 · 1812.96242004521 · 1991.54544060862 · 3328.32954924063 ·
25 6319.50600700604 · 1617.08143587061 · ... · 1401.38665678139 · 211.230377086175 · 5512.96874972635 · 905.30720963108 · 1606.9997081257 · -200.21108941412 · 282.282787140372 · 1907.40763094951 · -165.423791492452 ·
26 2871.85019645871 · 756.420740163849 · 1219.0298878222 · 2654.43583527845 · -534.0120203163249 · 2470.75207218967 · 2379.3502309867 · 1652.79132228454 · 1859.33177289075 ·
27 2371.96380531119 · 2688.406185803 · 14029.3072435639 · 10520.8017668873 · 14382.8917536474 · 10966.7151599875 · 5014.02879804059 · 6442.86197157372 · 4496.32603987551 · 3977.618011083 · 8496.3761331036 ·
28 5528.85818010868 · 5152.20846853366 · 3439.6040882518 · 1219.4427681081 · 12178.4427681081 · 4845.14822524977 · 4286.24275060326 · 1046.723476170708 · 1045.46561621815 · 1084.8768522354 · 4344.21759202625 · 1083.32855866281 ·
29 4489.48620648048 · 662.966005103606 · 2412.08724861677 · -337.531897633703 · 2385.78133194427 · 4976.5227143229 · 3335.86550372904 · -155.294419014334 · 2593.585309204334 ·
30 -180.235983768056 · 3088.90360026438 · 3957.33060681546 · 19524.0423334083 · 4159.91737554586 · 701.46459416793 · 8220.59906572288 · 10557.445479329 · 1488.8106533399 · 15230.543225621 · 2627.61869952451 ·
31 888.85171317634 · -60.8152721310253 · 14711.8316080412 · 5203.26622058384 · 4567.85649246882 · 16560.6277320492 · 8295.6418402986 · 1854.8986515611 · 7083.1605387531 · 1067.2666186641 · 4279.7497617058 ·
32 1324.36079187922 · 971.76747094036 · 174.504848722436 · 15226.1983054792 · 17672.8789568185 · 3594.960921061 · 8390.8444438924 · 18692.3706268832 · 3726.64632193648 · 2181.95823609631 · 9020.11632741929 ·
33 4992.8719710266 · 6662.14761000228 · 3777.1063094431 · 1927.85587461856 · -5.92081720892111 · 743.2207721564 · 10741.1800749465 · 1606.71051470696 · 10196.0774934666 · 6112.90493111247 · 5606.44776594194 ·
34 4742.15925919449 · 2621.97515141553 · 2673.34645293456 · 5326.59789040273 · 4574.96446398733 · 6632.86755747059 · 9796.63751569144 · 4527.97212974177 · 11350.1028996748 · 4723.67657514433 · 10038.8838594692 ·
35 2398.2797200203 · 483.121926367908 · 4530.72486557457 · 837.66740208941 · 8538.50188691285 · 12958.0582920047 · 1376.51129768868 · 2861.12843874583 · 2314.27149756348 · 912.0257915163041 · 7565.43344307537 ·
36 3936.73277172309 · 2413.39824835584 · 1026.47501065884 · 3047.96756049347 · 2207.47345879422 · 2421.1248136544 · 6793.21884765188 · 8158.70363801863 · 9146.23101766974 · 10723.4319256673 ·
37 2343.84724279661 · 790.04578696312 · 4838.08265267316 · 2413.83595153003 · 15537.3417583405 · 4848.29566117498 · -190.11454774273 · 1965.48092049776 · 465.21968857757 · 379.5981414285 · 3189.83240547675 ·
38 3555.53390149075 · 8034.8645320058 · 4932.2041096781 · 3231.45394882419 · 16913.801294608 · 1697.34743656757 · 982.66064301264 · -453.524807502577 · 1545.20732192982 · 811.71335452225 · 1445.25474793952 ·
39 1572.59681092731 · 3154.41729314318 · 1272.06943743717 · 8833.8716888124 · 11760.0640588486 · 2010.54165040329 · 1042.86457592954 · 235.6717320955072 · 4114.69539526113 · 509.830346913673 ·
40 2687.7472031583 · 956.041807038659 · 3955.9924391938 · 4511.0724888022 · 2021.6808748147 · 1664.0917041993 · 3241.51525021265 · 1344.02899771818 · 10605.7945031588 · 1714.9474073681 · 2079.77607485227 ·
41 5289.32086313067 · 2732.771911904495 · 499.494616780044 · 1167.78163298595 · 604.95328190718 · 1994.87192561236 · 2410.08539744206 · 1976.54024857554 · 678.498044753921 · 1378.16792475175 · 1679.55064179461 ·
42 1663.81827925231 · 2330.25366016752 · 481.810556227269 · 441.436432282926 · 4243.70200630799 · -123.20316209581 · 1071.26759868883 · 2732.09074364256 · 24.2182114644347 · 2592.03736172935 · 1712.82688752068 ·
43 3759.018180120303 · 2697.19844771165 · 3263.52094856016 · 2901.81286570455
```

\$coefficients

NULL

\$qgammas

\$coefficients

NULL

\$gammams

NULL

\$PRESSO

Apps: 1043066442.2424

\$PRESS

A matrix: 1 × 17 of type dbl

1 comps	2 comps	3 comps	4 comps	5 comps	6 comps	7 comps	8 comps	9 comps	10 comps	11 comps	12 comps	13 comps	14 comps	15 comps	16 comps	17 comps	
Apps	2416784138	1643495381	1462478268	1273898165	997760139	908050267	902532919	889400868	887550310	884723067	881210878	878661551	877937594	876717149	876431156	876408362	876437949

\$adj

A matrix: 1 × 17 of type dbl

1 comps	2 comps	3 comps	4 comps	5 comps	6 comps	7 comps	8 comps	9 comps	10 comps	11 comps	12 comps	13 comps	14 comps	15 comps	16 comps	17 comps	
Apps	3689126	2421610	2087662	1606026	1284484	1199941	1189753	1179679	1170930	1167448	1164582	1162699	1161976	1161680	1161616	1161616	1161616

\$segments

\$V1

469 · 76 · 423 · 251 · 321 · 314 · 331 · 280 · 246 · 511 · 199 · 183 · 266 · 539 · 397 · 569 · 102 · 501 · 463 · 224 · 317 · 345 · 608 · 245 · 322 · 376 · 310 · 464 · 453 · 302 · 621 · 416 · 577 · 609 · 489 · 594 · 176 · 25 · 395 · 494 · 557 · 335 · 410 ·

582 · 110 · 6 · 346 · 57 · 466 · 3 · 21 · 563 · 148 · 145 · 193 · 213 · 82 · 236 · 24 · 178 · 479 · 200 · 18 · 18

\$V2

70 · 579 · 249 · 333 · 214 · 407 · 353 · 46 · 598 · 242 · 443 · 446 · 325 · 533 · 448 · 434 · 104 · 618 · 601 · 522 · 123 · 284 · 480 · 12 · 165 · 26 · 341 · 390 · 543 · 391 · 202 · 263 · 360 · 13 · 34 · 503 · 573 · 499 · 15 · 477 · 572 · 316 · 157 · 197 ·

146 · 394 · 117 · 537 · 312 · 265 · 587 · 370 · 339 · 401 · 271 · 133 · 105 · 452 · 439 · 141 · 286 · 357 · 14

\$V3

276 · 532 · 59 · 364 · 597 · 505 · 328 · 27 · 546 · 223 · 191 · 551 · 261 · 139 · 185 · 329 · 124 · 383 · 343 · 270 · 351 · 576 · 510 · 371 · 275 · 173 · 382 · 589 · 50 · 11 · 211 · 179 · 330 · 201 · 516 · 393 · 490 · 290 · 571 · 488 · 313 · 109 · 388 · 367 ·

228 · 550 · 467 · 91 · 90 · 19 · 384 · 584 · 355 · 459 · 365 · 126 · 125 · 156 · 1 · 159 · 295 · 332

\$V4

465 · 256 · 581 · 175 · 447 · 115 · 17 · 396 · 255 · 526 · 30 · 413 · 512 · 320 · 350 · 491 · 490 · 283 · 85 · 362 · 622 · 205 · 406 · 44 · 545 · 568 · 363 · 229 · 68 · 216 · 549 · 192 · 235 · 417 · 361 · 260 · 400 · 239 · 586 · 88 · 529 · 221 · 155 · 298 ·

482 · 523 · 131 · 111 · 485 · 262 · 606 · 366 · 348 · 174 · 127 · 249 · 292 · 476 · 500 · 418 · 93 · 531

\$V5

75 · 602 · 137 · 234 · 521 · 294 · 380 · 49 · 457 · 232 · 580 · 495 · 619 · 101 · 287 · 496 · 129 · 42 · 300 · 304 · 555 · 438 · 60 · 356 · 103 · 10 · 147 · 293 · 471 · 519 · 592 · 16 · 540 · 441 · 308 · 167 · 427 · 398 · 342 · 535 · 28 · 528 · 305 · 481 ·

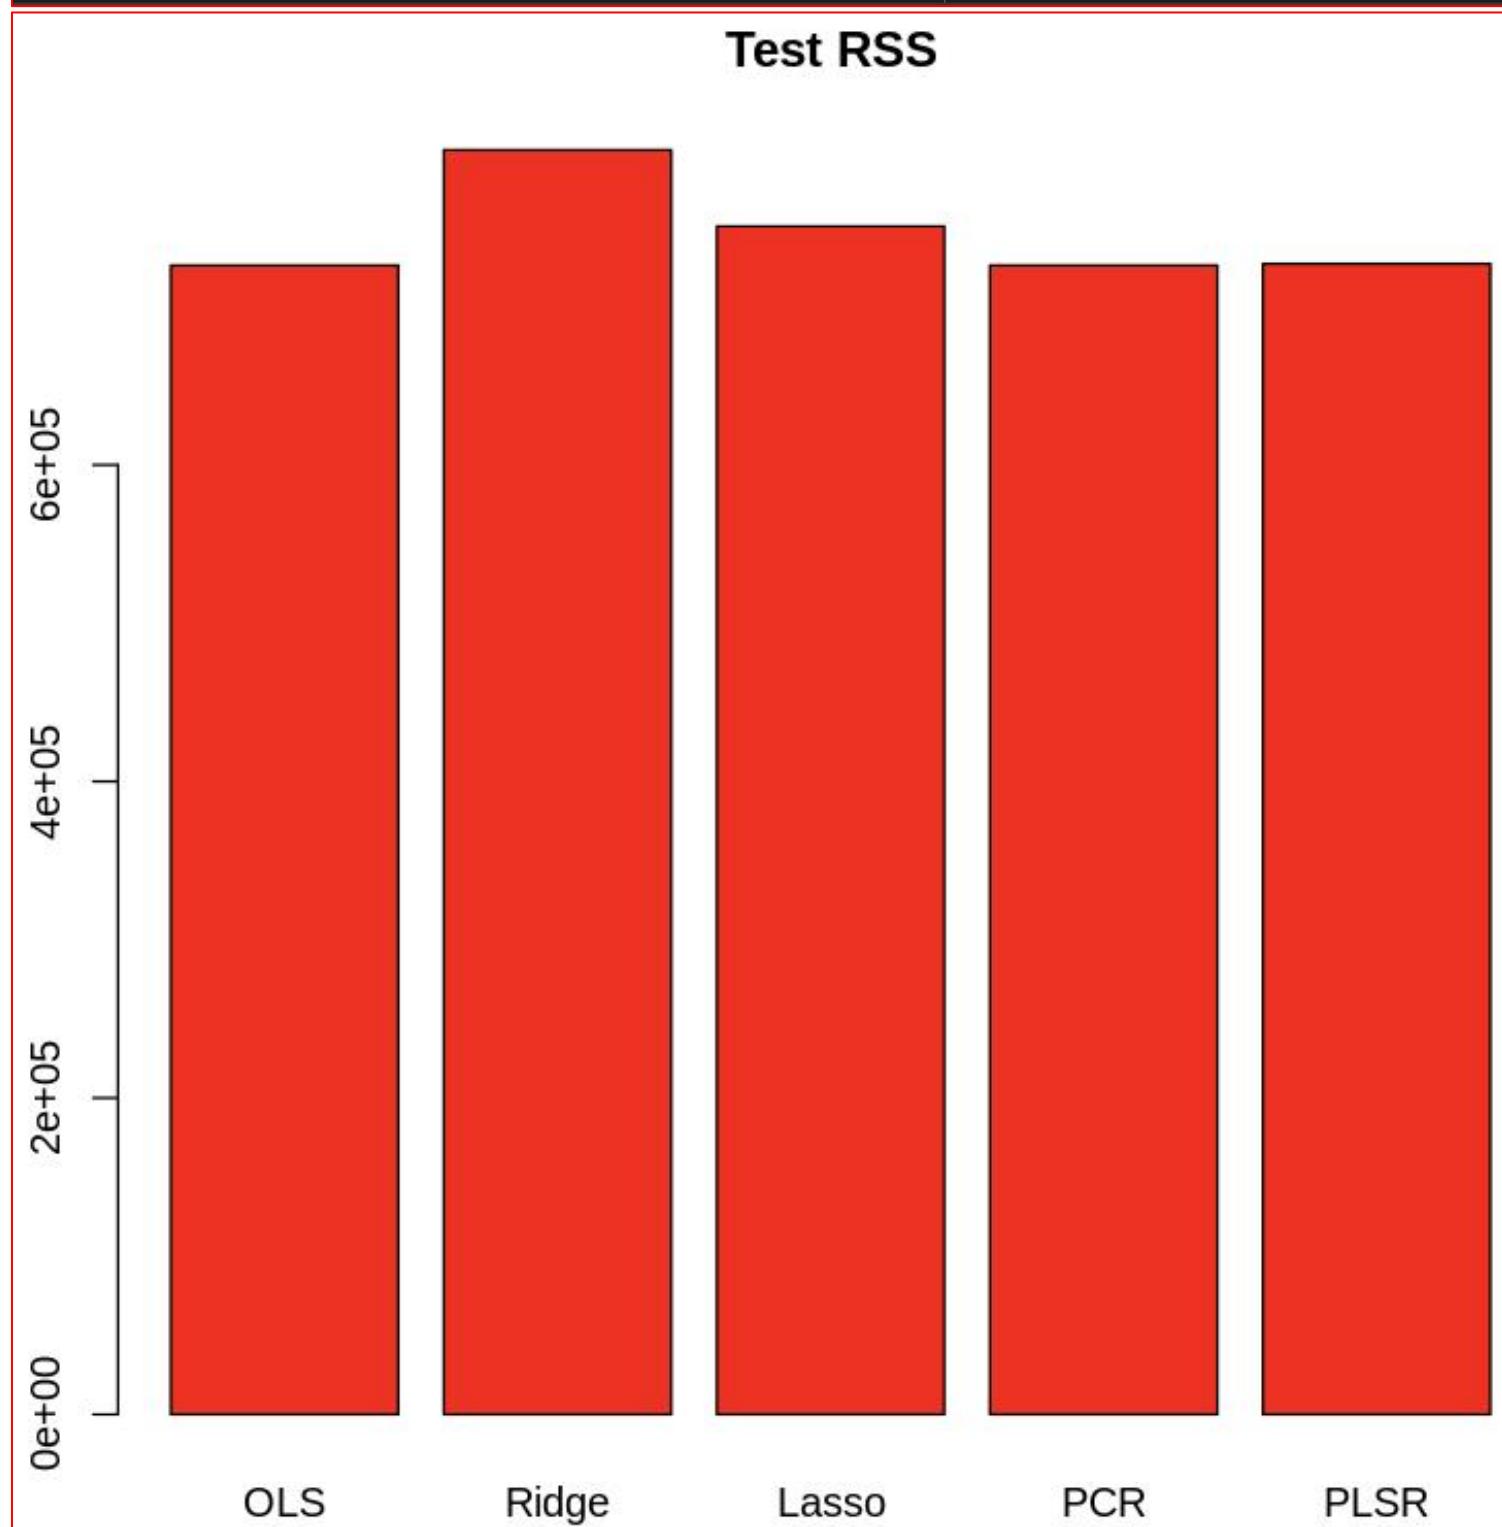
128 · 267 · 281 · 389 · 552 · 20 · 56 · 315 · 599 · 158 · 567 · 378 · 180 · 620 · 231 · 404

\$V6

296 · 132 · 198 · 272 · 67 · 542 · 408 · 583 · 354 · 37 · 181 · 51 · 241 · 604 · 507 · 187 · 402 · 451 · 190 · 89 · 392 · 233 · 419 · 422 · 94 · 368 · 403 · 442 · 493 · 254 · 169 · 47 · 437 · 386 · 98 · 277 · 258 · 1

(g) Comment on the results obtained. How accurately can we predict the number of college applications received? Is there much difference among the test errors resulting from these five approaches?

```
barplot(c(726062.2,799031.9,750801.5,726062,727135), col="red", names.arg=c("OLS","Ridge","Lasso","PCR","PLSR"),main="Test RSS")
```



Based on the given graph we can see that Ridge has the highest RSS (799031.9) as compared to others which is followed by Lasso (750801.5), then PLS (727135) and lastly Linear Model and PCR gave almost the same results for RSS with 726062 and 726062.2 respectively. Based on this we can Use any of the models except Ridge and Lasso.